

MEASURING AI'S IMPACT ON THE JOB SEARCH

Nicholas Qiu

Advised by Benjamin Golub

A thesis submitted for the Mathematical Methods in the Social Sciences
major

June 5, 2026

Abstract: AI-powered tools have reshaped hiring on both sides of the market, yet candidate behavior in this environment remains poorly understood. This paper examines how job-seeking students perceive the prevalence and ethical acceptability of AI-assisted cheating in recruitment, and whether beliefs about peer behavior shape their own willingness to engage in it. Drawing on a survey of 249 U.S. undergraduate, graduate, and early-career respondents, I test five hypotheses derived from focus theory, moral disengagement theory, and deterrence theory. Candidates' responses are consistent with an implicit hierarchy: preparation-stage AI use attracts far greater tolerance than live substitution during interviews. Perceived peer cheating shifts that hierarchy meaningfully, while process unfairness expands ethical tolerance for questionable behaviors selectively rather than wholesale. Contrary to classical deterrence theory, beliefs about the detectability of cheating showed no systematic relationship to willingness to cheat; penalty magnitude appears to matter more. AI agents replicated stable social norms well but failed to simulate the contextual sensitivity driving the most consequential variation in human responses — a finding with direct implications for hiring platforms that rely on AI-generated behavioral predictions.

1 Introduction

AI-powered tools have reshaped how both candidates and employers navigate the hiring process, reaching from resume screening and applicant tracking to interview preparation and automated application generation. What began as an employer-side efficiency tool has rapidly become a resource that candidates themselves deploy, raising a new and underexplored set of questions about conduct, fairness, and norms. By lowering the barrier to generating polished applications and simulating interview performance, AI has created new opportunities for candidates to misrepresent their abilities, while making it far less clear where legitimate assistance ends and dishonest conduct begins.

Yet despite the scale of this shift, the field knows remarkably little about how candidates themselves reason about it. In this paper, I aim to answer the question: How do job-seeking students perceive the prevalence and ethical acceptability of AI-assisted cheating in recruitment, and to what extent do perceptions of peer behavior shape individual willingness to engage in it?

The number of bachelor’s degree graduates among adults increased 715.2% between 1940 and 2020, and the job market has grown more competitive in step (Hanson, 2025). The proliferation of large language models has further accelerated this dynamic, enabling candidates to generate cover letters, tailor resumes, and submit applications at a scale and speed that was previously impossible. Existing research has examined AI’s role in employer-side screening, notably documenting concerns about algorithmic bias and the opacity of automated decision systems (Raghavan et al., 2020), but has paid considerably less attention to the candidate side of the dynamic. Prior work has not systematically examined how candidates define the ethical boundaries of AI use in the application process, or whether beliefs about peer behavior shift their own willingness to engage in conduct they might otherwise consider inappropriate. This matters because descriptive norm perception, or beliefs about what peers are doing, is a well-established driver of individual conduct (Cialdini et al., 1991): when individuals believe misconduct is widespread, they are more likely to engage in it themselves, regardless of their personal ethical assessments. The causal logic runs from perceived prevalence to norm inference to behavioral willingness: if candidates believe AI-assisted cheating is common, that perception may itself license conduct they would otherwise decline. If this dynamic is already operating among job-seeking students, it has direct implications for how recruiters should design hiring processes and how other institutions, such as universities and policymakers, should respond.

To answer these questions, I administered a survey to a sample of U.S. undergraduate, graduate, and recently graduated students who are actively searching for jobs or internships. The survey was designed to capture both respondents’ own reported behaviors and their estimates of peer behavior, a dual-measurement approach motivated by research on social desirability bias and pluralistic ignorance. Because candidates are unlikely to self-report ethically questionable conduct at its true rate, asking respondents what they believe their peers are doing provides a complementary signal. Prior work in social psychology has shown that peer estimates, while imperfect, are systematically related to actual group-level norms and can be used alongside self-reports to infer underlying population distributions (Cialdini et al., 1991). Respondents who completed an optional extended module were

awarded additional raffle entries as an incentive to do so.

As a methodological complement, I also deployed a cohort of AI agents that completed the same survey instrument. The comparison is motivated by a conceptual concern: AI tools are increasingly used in hiring to anticipate or model how candidates will behave, yet it remains untested whether those tools accurately capture the reasoning real candidates bring to ethically charged situations. Where AI-predicted responses diverge from what human candidates actually report, those gaps reveal the limits of what large language models understand about how people navigate normative uncertainty. This is especially consequential when norms are shifting rapidly and training data may lag behind evolving human attitudes, and where the cost of systematic misrepresentation is borne by the candidates being modeled.

Prior research on AI and the labor market clusters into three bodies of work, each of which illuminates part of the problem this paper addresses while leaving another part unexamined.

The first concerns AI’s structural effects on labor demand and competition. Chen et al. (2024) found that AI-powered automation decreased labor demand by approximately 17% for automation-prone jobs, with relevant firms posting roughly 95 fewer positions per quarter. Critically, this contraction concentrated among roles most susceptible to automation, suggesting that the candidates most likely to turn to AI assistance in their job search are also those facing the sharpest deterioration in opportunity. Rising competition for a shrinking pool of positions sharpens the incentives to use every available tool, including AI, to gain an edge, which in turn raises the question of where candidates draw the line between optimization and misrepresentation.

The second body of work concerns employer-side AI adoption and the tensions it has produced. Horodyski (2023) found that 73% of recruiters were willing to use AI tools in the recruiting process and 70% agreed that AI increases efficiency, indicating broad institutional acceptance of AI as a screening instrument. Yet more than 70% of the same recruiters expressed concern about algorithmic bias and the reduction of human judgment in candidate evaluation. This internal tension between adoption and reservation is significant but incompletely understood: because the study relies on cross-sectional self-reports, it cannot establish whether fairness concerns translate into changes in how recruiters actually weight AI-generated outputs. Candidates, for their part, share the skepticism. Zhang and Yench (2022) found that survey respondents do not consider hiring algorithms fair or effective, with women reporting more negative assessments overall. This suggests meaningful demographic heterogeneity in how candidates experience AI-mediated hiring, which motivates my attention to subgroup differences in perception. What neither study addresses, however, is how candidates reason about their own AI use rather than the employer’s.

The third and thinnest body of work concerns candidate-side behavior during the application process itself. Galdin and Silbert (2025) offer the most direct evidence available: in a controlled setting where candidates were explicitly permitted to use AI, average application time fell from 1.47 to 1.08 minutes, and AI use depressed the quality of job market signaling. If a competitive application can be generated in just over a minute, the informational value of application materials as a signal of candidate quality is substantially diminished. Employers can no longer

treat polished applications as reliable indicators of underlying skill, and candidates who invest genuine effort may find themselves indistinguishable from those who did not. Crucially, however, Galdin and Silbert (2025)’s study was conducted under conditions of explicit permission. It does not address what candidates do when AI use is neither sanctioned nor prohibited, or how they evaluate the ethics of their own and others’ conduct in that ambiguous space.

Read alongside one another, these studies reveal a consistent asymmetry: candidates distrust employer-side AI tools as unfair and opaque, yet are themselves using candidate-side AI in ways that may equally distort the hiring process. What neither line of research has examined is how candidates reason about that asymmetry. Specifically, where do they draw the line between acceptable optimization and dishonest conduct, and does the belief that peers are crossing that line shift their own behavior? Answering these questions requires studying the candidates navigating this uncertainty directly. Undergraduate and graduate students represent a particularly informative population for this inquiry. They are entering the most AI-saturated hiring environment in recent memory, without the professional anchors that more experienced workers may have: established norms, institutional guidance, and accumulated experience. They are also at the life stage where behavioral norms are most actively formed, making their perceptions and choices especially consequential for how AI use in hiring develops over the coming decade.

The paper proceeds as follows. Section 2 describes the survey instrument, sample characteristics, and the design of the agentic study. Section 3 presents the theoretical framework and five hypotheses, along with the analytical approach. Section 4 reports findings organized around those hypotheses. Section 5 discusses shortcomings and limitations of the research. Section 6 presents theoretical and practical implications and concludes with directions for future research. Section 7 concludes the paper with a summary of key findings and their significance for the future of AI in hiring.

2 Survey and Sample

2.1 Data Collection and Final Sample

Survey-based methods are well suited to this study because the core constructs of interest, such as attitudes toward AI use, perceived peer behavior, and self-reported recruitment conduct, cannot be recovered from employer records, platform logs, or administrative data. Many of the behaviors under study also carry implicit social stakes, as applicants are unlikely to disclose AI-assisted application practices in professional settings. Surveys make it possible to elicit candid self-reports at scale under conditions of anonymity, enabling the detection of attitudinal patterns and group-level differences that would otherwise remain invisible.

The survey was administered between December 2025 and May 2026, collecting complete responses from 249 U.S. undergraduate students, graduate students, and early-career professionals. Administration combined online

distribution through Qualtrics with in-person recruitment at campus events and study spaces, with the latter targeting populations with lower rates of response to cold online outreach. Respondents were recruited primarily through Northwestern University, the University of Chicago, and several other institutions in the greater Chicago area. A smaller share came from institutions outside Chicago, reached through student organizations and professional networks. The responses, in addition to the analysis code and survey instrument, are available in a public repository on GitHub.¹

The sample was intentionally concentrated in STEM fields and related quantitative disciplines, including Computer Science, Data Science, and finance. These populations were prioritized because prior research indicates that AI-assisted tools have been adopted earlier and more extensively in technical fields than in humanities or social science disciplines (Contractor & Reyes, 2025). The technical and credential-intensive structure of STEM and finance recruiting, in which interviews frequently involve coding exercises, quantitative problem sets, and case assessments, also creates direct and well-defined opportunities for AI assistance, making these groups an especially relevant population for this study. Respondents from other fields were also recruited, though at lower rates. This under-representation of non-STEM perspectives reflects a genuine limitation of the convenience sampling approach, as recruitment networks and in-person access were concentrated in technical programs. This constraint is addressed further in the limitations section.

Social desirability bias is a recognized limitation of self-report data: respondents may underreport behaviors they perceive as unacceptable and overreport those they consider normative. The primary mitigation strategy draws on the social psychology literature on pluralistic ignorance, which shows that personal beliefs and perceived peer beliefs often diverge systematically, and that comparing the two can reveal underlying group norms more reliably than self-reports alone (Baugh et al., 2021). The survey therefore included parallel items asking respondents to report both their own attitudes and conduct and what they believe to be true of their peers. Respondents who underreport their own behavior may nonetheless report accurately what they believe peers do; this gap, if it exists, is treated as evidence of social desirability distortion rather than genuine disagreement. Peer-belief items therefore serve as an internal check on distortion rather than a redundant measure. To further reduce normative pressure, questions were framed descriptively rather than evaluatively, asking respondents to characterize behavior rather than render judgment on it.

Respondents were contacted through one of three channels: a direct survey link, a QR code distributed in person or online, or an in-person request. Upon opening the survey, participants were presented with a consent page explaining that the study was conducted for academic research purposes by a nonpartisan researcher, that participation was voluntary, and that all responses would be anonymous. Demographic questions were placed at the end of the survey rather than the beginning, both to maximize completion of the primary items and to avoid priming effects that might shape how respondents interpreted earlier questions. As a participation incentive, respondents who completed the survey were entered into one of three raffles for a \$50 Amazon gift card. Respondents who completed an optional supplementary section were awarded two additional raffle entries. Table 1 summarizes the

¹<https://github.com/nickqiu21/ai-job-search-study/> contains the full dataset, code for all analyses, and the survey instrument.

demographic characteristics of the sample. Note that gender and age were not collected to preserve anonymity, as the combination of field of study, academic level, and interview experience was sufficient to create unique respondent profiles in some cases. The sample is therefore best described as a young, early-career population with a strong representation of technical fields, which is consistent with the study’s focus on candidates navigating AI-mediated hiring processes.

Prior to full deployment, the survey underwent IRB review. Instrument development included two phases of pilot testing. In the first phase, a small cohort of target-population participants (n=11) completed the survey and provided feedback on question clarity, response option interpretability, and the overall length and flow of the instrument. In the second phase, synthetic responses were generated using (ExpectedParrot), a platform for simulating survey responses using large language model agents, drawing profiles from the target population and varying them across field of study, academic level, and reported AI tool usage. This simulation served two purposes: it provided a baseline distribution of expected responses against which the human pilot data could be compared, and it surfaced potential sources of ambiguity in question phrasing where AI-generated responses diverged sharply from human responses. The synthetic responses were not used as analytical data; they functioned as an instrument validation tool prior to human data collection.

2.2 Survey Structure

The full questionnaire is available in Appendix A.1. Figure 1 shows a full schematic of the full block structure; the following describes the purpose and measurement logic of each block.

Block 1 presented general questions about routine academic and professional behaviors, including frequency-of-use items designed to establish behavioral baselines and reduce priming effects before introducing questions directly related to AI or recruitment. The block also included an attention check item to support data quality screening.

Block 2 asked respondents to rank eight potential motivations for using unauthorized assistance during job interviews, and included an open-ended item for elaboration. The ranking format was chosen over Likert-style agreement items to force discrimination among motivations that respondents might otherwise endorse uniformly, and to allow respondents to surface considerations not anticipated by the fixed response options.

Block 3 used an anonymous coding protocol in which respondents were assigned a personal code to answer yes/no items about whether they were aware of others having used external assistance during recruitment or academic assessments. Prior research establishes that misreporting on sensitive topics is largely situational, shaped by the design features of the survey rather than by fixed respondent characteristics (Tourangeau & Yan, 2007). Indirect elicitation methods, where respondents report on third-party conduct rather than their own, are a recognized strategy for reducing this pressure by lowering the perceived personal stakes of disclosure (Stantcheva, 2023). The block also included a set of calibration check items, each posing a question with an objectively verifiable answer, to screen for respondents who were not engaging with the instrument in good faith.

Block 4 presented hypothetical scenarios in which an interview candidate learns that competing candidates are using AI assistance, eliciting respondents' judgments about appropriate conduct under conditions of perceived norm uncertainty. Two conditions were presented within-subject: a high-competition condition in which AI use among competitors is described as widespread, and an alternative condition in which it is not. Scenario-based items were used because direct attitude questions on a sensitive topic risk socially desirable responding; placing judgments in a third-person context reduces that pressure while still capturing the respondent's own normative reasoning.

Block 5 collected demographic information, including academic year, field of study, and prior interview experience as a job applicant. Demographics were positioned last to avoid priming respondents with identity-relevant categories before the substantive items.

An optional Block 6 assessed respondents' trust in the hiring process more broadly and solicited open-ended feedback. This block was included for exploratory purposes given limited prior literature on this population's institutional attitudes.

Upon completion of the survey, respondents that opted into the prize raffle were directed to a separate form to enter the prize raffle at the conclusion of the survey.

Table 1: Sample Characteristics (N = 249)

Characteristic	<i>n</i>	%
<i>Academic Level</i>		
Freshman	26	10.4
Sophomore	46	18.5
Junior	75	30.1
Senior	57	22.9
Graduate Student	29	11.6
New Grad (0–1 YOE)	10	4.0
Experienced (2+ YOE)	3	1.2
Other	3	1.2
<i>Field of Study^a</i>		
Economics	58	23.3
Computer Science	37	14.9
Data Science / AI / ML	23	9.2
Mathematics or Statistics	22	8.8
Mechanical Engineering	13	5.2
Biology	12	4.8
Other	84	33.7
<i>Desired Industry^a</i>		
Software Engineering / AI / Data Science	52	20.9
Management Consulting	41	16.5
Healthcare (clinical & research)	31	12.4
Finance (all tracks)	29	11.6
Other	96	38.6
<i>Interview Experience</i>		
0–1 interviews	38	15.3
2–6 interviews	113	45.4
7–10 interviews	45	18.1
11–20 interviews	32	12.9
21+ interviews	21	8.4
<i>Primary Interview Format</i>		
Online	174	69.9
Mixed	56	22.5
In-person	19	7.6

^a Full category distributions appear in Appendix A.3. Tail categories are collapsed into “Other” for readability.




Figure 1: Survey block schematic. Blocks are presented in fixed sequence. Color encodes functional role: purple for peer perception, teal for disclosure, coral for scenario judgment, and navy for structural framing. The optional block (dashed border) is reached via the branching node; participants who decline are routed directly to survey completion. Question numbers are shown for reference to the full instrument in Appendix A.1.

2.3 Agentic Study Design

To complement the human survey, a cohort of 150 AI agents was administered the identical survey instrument via [ExpectedParrot](#). The rationale for this parallel study is both conceptual and analytical: where AI-predicted responses diverge from what human candidates actually report, those gaps reveal the limits of what current large language models understand about how people navigate ethically ambiguous situations under social pressure. The agent cohort is treated throughout as a distributional baseline, or a record of what a rich predictive model produces on the same instrument that human respondents completed.

Each LLM agent was instantiated using the default ChatGPT 5.2 model and assigned a demographic profile drawn from the target population. Profiles varied across six attributes: age, gender, education level, years of professional experience, industry, and employment status. Additional profile characteristics were recorded on a 1 to 5 scale, including self-reported AI usage frequency and three baseline attitudinal measures: trust in AI tools, perceived fairness of AI-assisted hiring, and self-assessed confidence in obtaining an interview. These attitudinal attributes were included to introduce heterogeneity in the agent cohort beyond demographic characteristics alone, allowing for within-cohort variation in responses that might otherwise collapse to a single modal answer.

The resulting cohort of 150 agents spans an age range of 18 to 27, with a mean age of 21.7 years, closely matching the undergraduate and early-career graduate profile of the human sample. Educational attainment is similarly aligned: most agents hold bachelor’s degrees in progress, with smaller proportions holding completed bachelor’s or master’s degrees (see Appendix A.2 for the full distribution). Six industries are represented, selected to reflect the occupational diversity of the human sample while concentrating on fields where AI-assisted hiring is most prevalent: Technology, Healthcare, Marketing, Finance, Education, and Consulting. Employment status varies across actively applying (79 agents, 52.7%), passively looking (36, 24.0%), and already employed (35, 23.3%). AI usage frequency skews toward regular use, with 107 of 150 agents (71.3%) reporting using AI tools “Often” or “Very Often.” Mean trust in AI tools was 3.60 on the 1-5 scale ($SD = 1.04$), perceived fairness of AI hiring was 3.11 ($SD = 1.06$), and confidence in obtaining an interview was 3.47 ($SD = 0.92$). The full distribution of agent profiles is available in the [GitHub repository](#) accompanying this paper.

Agents were prompted with minimal instruction by design: each was given its assigned demographic and attitudinal profile but received no guidance on how to respond to specific survey items, nor any direction to simulate particular ethical dispositions. This was intended to elicit responses that reflect the model’s own priors about how a person with those characteristics would reason, rather than responses shaped by researcher expectation. The tradeoff this creates is acknowledged directly: because agents received no instruction to simulate psychological phenomena such as response fatigue, mood variation, or intra-survey inconsistency, their responses are likely more uniform and internally consistent than those of human respondents. This constraint is discussed further in the limitations discussion in Section 5.

Systematic divergence between the two distributions will be of interest as a measure of “incremental” information

in the survey sample, examined alongside human findings in each hypothesis subsection of Section 4.

3 Analytical Strategy

3.1 Theoretical Framework

A recurring finding across social psychology is that individuals do not evaluate ethical conduct in isolation; their judgments are shaped by perceived norms, social context, and the institutional environment in which decisions are made. The growing integration of AI tools into hiring has created a specific and underexplored instance of this dynamic, one in which the boundaries of acceptable conduct are unsettled, competitive pressure is high, and the behavior of peers is difficult to observe directly. This section develops a theoretical framework for understanding how candidates navigate that environment. It draws primarily on social psychology, which offers a set of well-established mechanisms for explaining how perception, social context, and accumulated experience shape behavior in precisely these conditions. Four of those mechanisms—descriptive norm formation, moral disengagement, deterrence, and social learning—are examined in turn, before the section closes with a discussion of what the framework implies for the study’s hypotheses and for the AI agent comparison that accompanies them.

The first and most proximate mechanism is descriptive norm formation, as theorized by Cialdini et al. (1991) in their focus theory of normative conduct. Focus theory holds that individuals calibrate their behavior not primarily against abstract moral principles but against their perception of what peers in the same situation are doing. When a candidate believes that many competitors are using AI assistance, that belief functions as a descriptive norm, signaling that the behavior is common, and therefore implicitly permissible, regardless of whether the candidate personally endorses it. Critically, the operative variable is perceived prevalence, not actual prevalence. A candidate who overestimates how many peers are cheating will respond to that inflated estimate as if it were accurate, because the norm is defined by what the candidate believes rather than what is objectively true. This has a direct implication for the present study: the distribution of peer cheating estimates, not just whether candidates cheat, is a primary object of analysis.

Descriptive norms establish the social reference point, but they do not fully explain why some candidates act on that reference point while others do not. A second mechanism, moral disengagement, accounts for this variation. Bandura (1999) argues that individuals do not simply apply a fixed ethical standard to every situation; rather, they are capable of suspending ethical self-regulation when conduct can be reframed as a justified response to a broader wrong. In the context of AI-assisted hiring, the relevant reframe is fairness restoration: if the hiring process is perceived as systematically unfair, candidates who might otherwise decline to cheat can reconstruct the same conduct as a proportionate response rather than a violation. The moral cost does not disappear but is offset by the perceived injustice of the system. This mechanism predicts that candidates who view the process as unfair will not simply be more cynical in a diffuse way; they will be more willing to engage in specific behaviors they would

otherwise consider unacceptable, and the size of that shift will vary across behaviors depending on how directly the reframe applies.

These two mechanisms together describe how the moral threshold for cheating shifts before the candidate acts. A third mechanism, deterrence, shifts the frame from moral to instrumental: candidates do not only ask whether cheating is wrong, but whether it is worth it. Classical deterrence theory (Becker, 1968) holds that behavior is restrained by the expected cost of deviation, which is a function of both the probability of detection and the severity of the consequence if caught. Where detection is perceived as unlikely and consequences as mild or inconsistently applied, the expected cost of cheating falls, and the deterrence mechanism weakens. In competitive hiring environments where interviews are conducted remotely and proctoring is limited, candidates have reasonable grounds for believing that AI assistance is difficult to identify. If they also believe that companies rarely act on suspicion even when it arises, the instrumental calculus that deterrence assumes predicts higher stated willingness to cheat, independent of whatever attitudinal shifts descriptive norms and moral disengagement have already produced.

The three mechanisms above describe the conditions that make cheating more likely at a given moment. A fourth mechanism, social learning, explains how those conditions change with experience. Social learning theory holds that individuals update their behavioral norms through observation of what peers actually do and what consequences actually follow. Each additional hiring cycle gives a candidate more data: more observations of whether peers used AI assistance, more evidence about whether cheating was detected, and more exposure to the reframings that experienced candidates use to normalize conduct they once found uncomfortable. This study introduces the concept of ethical disillusionment to describe the gradual replacement of fixed ethical limits with situational reasoning as hiring experience accumulates. The concept raises the question of whether the effects of descriptive norms, moral disengagement, and deterrence will be more pronounced among experienced candidates than among first-time applicants.

These four mechanisms do not operate in isolation. A candidate who perceives widespread peer cheating, views the process as unfair, and doubts that detection is likely is exposed to three separate inputs, each of which the literature treats as independently sufficient to shift behavior. Together, they describe not an additive combination but a shared environment in which the moral threshold for cheating is depressed from multiple directions at once. Accumulated experience, moreover, shapes how each of those perceptions formed, making it less a fourth input than a conditioning factor on the other three. Whether these mechanisms reinforce one another in practice, and by how much, is an empirical question the analyses below are designed to examine.

This framework also motivates the fifth hypothesis, which concerns the AI agent cohort. What the four mechanisms share is their dependence on contingent, experiential inputs: social pressure that accumulates across cycles, institutional trust that shifts depending on what candidates observe, and competitive context that changes with each new hiring round. These are outcomes of an unfolding process rather than fixed traits readable from a demographic profile. It is genuinely unclear whether models trained on population-level patterns can replicate the variation that emerges from the market as it has actually played out, and in particular whether they can capture

the distributional surprises that experience tends to generate. This study treats that uncertainty as an object of inquiry in its own right.

3.2 Hypotheses

Taken together, these mechanisms produce five hypotheses that the empirical analysis is designed to evaluate:

H1 (Descriptive Norm Effect): Higher perceived prevalence of cheating among peers will be associated with greater personal willingness to engage in unauthorized assistance during interviews, consistent with focus theory’s prediction that descriptive norms are a primary driver of individual behavior (Cialdini et al., 1991).

H2 (Fairness and Moral Disengagement): Lower perceived fairness of the hiring process will be associated with greater willingness to cheat. Candidates who view the process as unjust are expected to perceive a lower moral cost to dishonest conduct, as moral disengagement theory predicts that individuals are more likely to suspend ethical self-regulation when rule-breaking can be framed as a justified response to an unfair system (Bandura, 1999).

H3 (Deterrence): Lower confidence in companies’ ability to detect or act on cheating will be associated with greater willingness to engage in it, as reflected in both self-reports and social perceptions. When candidates believe consequences are unlikely, the expected cost of cheating falls, producing higher stated willingness.

H4 (Experience and Normalization): Greater prior interview experience will be associated with higher normalization of dishonest behavior. This prediction tests the ethical disillusionment concept, treating interview exposure as a proxy for the accumulated experiential inputs through which the three preceding mechanisms compound over time.

H5 (Human-AI Divergence): If willingness to cheat is shaped primarily by contingent, experiential inputs rather than stable demographic traits, AI agents reasoning from persona profiles may struggle to reproduce the distributional patterns observed in human respondents. This study examines whether meaningful divergence exists between agent and human responses, and if so, what those gaps reveal about the limits of profile-based reasoning in ethically ambiguous and socially contingent situations.

3.3 Data and Analytical Approach

Survey responses collected via Qualtrics and in-person administration were analyzed using Python (pandas, scipy, statsmodels) and Excel for descriptive tabulation. The analyses are organized around the five hypotheses above, with each analytical block designed to test a specific prediction rather than describe a general pattern.

To test H1, I examined the correlation between perceived peer prevalence of cheating and individual willingness to engage in it, assessing whether higher descriptive norm estimates predicted greater stated willingness independent

of personal ethical assessments. As a complement, I tested whether knowing at least one person who had cheated produced a similar association, treating this as a more concrete behavioral anchor for the norm perception variable. To formalize the peer influence mechanism, I applied a DeGroot (1974) social contagion framework, which models opinion updating through iterated peer observation. Each respondent is treated as a node whose willingness to cheat is a function of perceived cheating frequency among competitors for the same position, allowing estimation of the magnitude and convergence properties of peer influence on individual willingness.

To test H2, I analyzed the relationship between perceived fairness of the hiring process and willingness to cheat. This included both direct correlational analysis and a comparison between respondents who rated the process as unfair and those who did not, examining whether fairness perceptions shifted ethical evaluations of specific cheating behaviors as well as general willingness.

To test H3, I compared confidence in companies' ability to detect cheating against confidence in candidates' own ethical conduct, assessing whether perceived detection risk moderated willingness. Respondents reporting low detection confidence were expected to report higher willingness, consistent with a deterrence logic.

To test H4, I compared respondents with greater interview experience against less experienced respondents on reported cheating prevalence, personal willingness, and ethical evaluations of specific behaviors, using median experience as the split criterion. This analysis was designed to assess whether exposure to the hiring process produces the normalization pattern predicted by the ethical disillusionment concept, or whether prior experience does not predict the degree of normalization observed.

To test H5, I compared the distributions produced by human respondents against those generated by the AI agent study, examining both overall distributional differences and directional trends in the relationship between process confidence and cheating likelihood. Systematic underestimation by AI agents was defined as a statistically meaningful gap between model-predicted willingness and human-reported willingness, assessed using distributional comparison tests across the same survey instrument.

In addition to these hypothesis-driven analyses, I conducted two supplementary investigations. First, group-based analysis categorized respondents by willingness profile: unwilling under any circumstance, conditionally willing, indifferent, or generally willing. These groups were identified using k-means clustering on willingness-related items and then compared across industry, education level, and interview experience to identify whether particular subpopulations drive the aggregate patterns. Second, for respondents who completed the open-ended question, text and sentiment analysis identified recurring themes and justifications, using qualitative evidence to surface patterns and explanations that the quantitative data could not capture alone.

Incomplete survey responses were excluded from analysis; only respondents who completed all primary items were retained in the analytical sample. Because demographic questions were positioned at the end of the survey, incomplete responses preclude any systematic comparison of completers and non-completers on observable characteristics, and the direction of any resulting selection bias cannot be determined from the data.

4 Survey Analysis

4.1 H1: Descriptive Norm Effect

H1 predicts that higher perceived prevalence of peer cheating will be associated with greater individual willingness to engage in unauthorized assistance, consistent with focus theory’s account of descriptive norms as a primary behavioral driver (Cialdini et al., 1991). Four survey items bear directly on this hypothesis: Question 2, which elicited each respondent’s estimate of the share of job-seekers who have cheated on an interview; Questions 9 through 12, which asked whether respondents personally knew someone who had used AI or other unauthorized assistance during recruitment; Question 13a, which asked how likely the respondent would be to use unauthorized assistance if a high number of competing candidates were cheating; and Question 14a, which posed the same question under the counterfactual condition of zero peer cheating. Together, these items allow for both a correlational test of the norm-behavior relationship and a within-respondent comparison of how stated willingness shifts across competitive contexts. To formalize the peer influence mechanism underlying these patterns, the DeGroot model introduced in Section 3.2 is estimated as a network-level test of whether opinion convergence through industry peer updating predicts observed willingness.

Human respondents estimated, on average, that 36.7% of job-seekers they know have cheated on an interview. The distribution was notably right-skewed, with responses ranging as high as 95%, and a median of 33.0%. This stands in sharp contrast to the agent cohort, whose estimates clustered tightly around 12.7%, with responses limited almost entirely to 5, 10, and 15%. The gap is large enough to reflect qualitatively different reasoning processes, not merely different priors. AI agents appear to reason from idealized population norms, producing estimates consistent with low base rates of misconduct. Human respondents, by contrast, draw on their actual social networks, and in a sample concentrated in competitive STEM and finance programs, those networks plausibly contain real instances of observed or reported cheating. Whether or not the human estimates are accurate at the population level, they are the estimates that shape behavior: if a candidate believes that 20% of their peers have cheated, that belief functions as a descriptive norm regardless of its objective validity. Figure 2 plots the distribution of Question 2 estimates for both samples, and Figure 3 presents a scatterplot of peer cheating estimates against Question 13a willingness among human respondents; the correlation between the two (Spearman $r = 0.262$, $p < 0.001$) provides the most direct test of the norm-behavior relationship.




Figure 2: Distribution of peer cheating estimates (Q2) among human respondents (blue) and AI agents (red). Human estimates are right-skewed, with a mean of 36.7% and a median of 33.0%, while agent estimates cluster tightly around 12.7%. The 24.1 percentage-point gap suggests that human respondents draw on actual social networks, while agents reason from idealized population priors.




Figure 3: Peer cheating estimates (Q2) plotted against willingness to cheat under high peer cheating conditions (Q13a) among human respondents. A Spearman correlation of ($r = 0.262$, $p < 0.001$) confirms a positive relationship between perceived peer prevalence and stated willingness, consistent with the descriptive norm hypothesis.

The anonymous disclosure items provide a more concrete behavioral anchor for norm perception than the abstract percentage estimate of Question 2. A respondent who personally knows someone who cheated has direct, first-hand evidence of peer misconduct, and that evidence is likely to carry greater normative weight than an estimated population rate. Among human respondents, 89.2% reported knowing someone who had used AI to complete a homework assignment when its usage was explicitly prohibited (Q9), 37.3% knew someone who had used AI to generate interview answers in advance without authorization (Q10), 26.9% knew someone who had used a search engine or peer assistance live during an interview (Q11), and 26.1% knew someone who had used AI to generate responses live during a virtual interview (Q12). By contrast, only 3 out of 150 agents reported knowing someone who had used AI for unauthorized academic work, and no agent reported knowing anyone who had engaged in any of the interview-specific cheating behaviors. This gap reinforces the pattern observed in Question 2: agents systematically underestimate the prevalence of cheating within the social networks of real candidates, and that underestimation is sharpest for the behaviors most directly relevant to recruitment.

Cross-tabulating the disclosure items against Question 13a willingness reveals that respondents who reported knowing at least one person who had engaged in any of the four behaviors were 2.56x more likely to report being “Somewhat Likely” or “Extremely Likely” to cheat under high peer cheating conditions than those who reported no such knowledge. This association is consistent with the descriptive norm mechanism: personal exposure to cheating behavior among peers appears to shift the perceived normality of that behavior in ways that aggregate percentage estimates alone may not capture. Cross-sectional data cannot determine whether this reflects genuine norm updating through social observation or a selection effect in which candidates already more tolerant of cheating self-select into networks where cheating is more visible. The pattern is consistent with both accounts, and the distinction matters for interpretation.

When told that a high number of competing candidates were cheating, 55.0% of human respondents reported being “Somewhat Likely” or “Extremely Likely” to use unauthorized assistance themselves. This compares to 0% of agents, with 140 out of 150 reporting “Somewhat Unlikely” and the remaining 10 reporting “Extremely Unlikely.” The within-respondent shift is equally telling: when the competitive condition was reversed and respondents were told that no candidates were cheating, human willingness dropped sharply, with 62.2% now reporting “Extremely Unlikely” compared to 16.1% under the high-cheating condition. Agents showed no analogous sensitivity, with all 150 reporting “Extremely Unlikely” under the no-cheating condition, essentially unchanged from their already-low baseline. This flatness suggests that current large language models apply a consistent ethical framework regardless of situational context, while human respondents update their stated willingness in response to perceived competitive pressure. That difference is precisely what descriptive norm theory predicts, as illustrated in Figure 4.

A DeGroot social influence model was estimated to test whether a formal peer averaging process could account for the magnitude of the willingness shift from Q14a to Q13a. Using industry peer groups and Q2 cheating estimates as influence weights, the model produced converged willingness estimates that fell systematically below observed Q13a means across all industry groups, with an overall MAE of 2.055 and a near-zero Spearman correlation between predicted and observed group means ($r = -0.005$, $p = 0.932$). The model fails both in absolute terms and in




Figure 4: Willingness to cheat under high peer cheating conditions (Q13a) and no peer cheating conditions (Q14a) among human respondents and AI agents. The sharp within-respondent shift among human respondents, absent among agents, reflects the context-sensitivity that descriptive norm theory predicts.

recovering the rank ordering of industries by observed willingness, most likely because the Q14a baseline already embeds some implicit social reference and because the norm effect may operate through threshold dynamics that weighted averaging cannot reproduce. Full model specification and fit statistics are reported in Appendix A.4. The correlational evidence, however, does not require the model to succeed: the Spearman correlation between Q2 estimates and Q13a willingness ($r = 0.262$, $p < 0.001$), the disclosure-based multiplier, and the within-respondent scenario shift together constitute convergent evidence that the descriptive norm mechanism is operating, even if its precise functional form remains unresolved.

Taken together, the correlational, disclosure-based, scenario-based, and model-based evidence point toward the same conclusion: perceived peer cheating is positively associated with individual willingness to cheat, and this relationship is consistent with the descriptive norm mechanism that focus theory identifies. The DeGroot formalization of that mechanism, however, does not account for the magnitude of the shift from baseline to social-pressure willingness, suggesting that either the initial opinion measurement is biased, the averaging dynamics are the wrong functional form, or both. What remains unresolved is the causal direction: whether norm perception is genuinely driving willingness or whether both reflect a common underlying disposition toward conditional ethics. The experimental design needed to settle that question is discussed in Section 6.

4.2 H2: Fairness and Moral Disengagement

H2 predicts that lower perceived fairness of the hiring process will be associated with greater willingness to cheat, on the grounds that candidates who view the process as unjust face a lower moral cost to dishonest conduct. This follows from Bandura (1999)’s theory of moral disengagement, which holds that individuals are more likely to suspend ethical self-regulation when rule-breaking can be framed as a justified response to an unfair system. Three survey items bear on this hypothesis: Question 21, which elicited respondents’ overall fairness assessments of the modern job interview process; Question 3, which collected acceptability ratings for six specific cheating behaviors; and Question 6, which invited open-ended elaboration on why candidates cheat. Cross-tabulating fairness ratings against acceptability scores and drawing on qualitative themes from Question 6 allows for a test of whether perceived unfairness shifts ethical evaluations at the level of specific behaviors, not just general disposition.

On the fairness question, human respondents were moderately but not strongly skeptical. 43.1% rated the process as “Somewhat Unfair,” while 15.3% rated it “Neither Fair nor Unfair” and 23.4% “Somewhat Fair”. Notably, few respondents chose either extreme: only 16.8% reported the process as “Extremely Unfair” and 1.5% as “Extremely Fair,” suggesting that strong opinions in either direction are uncommon. This tempered skepticism is broadly consistent with Zhang and Yench (2022), who found that candidates tend to view AI-mediated hiring as neither fully fair nor fully illegitimate. The agent distribution was similarly split, with 59 out of 150 agents rating the process “Somewhat Unfair” and 91 rating it “Somewhat Fair,” and no agent selecting an extreme response. The convergence between human and agent fairness ratings is one of the closer alignments between the two samples in this study, though as the cross-tabulation below shows, the behavioral implications that humans draw from those fairness assessments diverge considerably from what agents predict. Figure 5 displays the distribution of fairness ratings across both samples.




Figure 5: Distribution of fairness ratings for the modern job interview process among human respondents and AI agents. Human respondents centered around ‘Somewhat Unfair,’ with a notable share selecting ‘Extremely Unfair’ (16.8%), while agents split almost entirely between ‘Somewhat Unfair’ and ‘Somewhat Fair’ with no extreme responses. Despite this partial convergence in the moderate range, the behavioral implications that humans and agents draw from similar fairness ratings diverge considerably, as subsequent analysis shows.

The acceptability ratings from Question 3 reveal a clear internal hierarchy among human respondents. Having someone else provide live assistance during an interview was rated least acceptable, with agents averaging 1.03 and humans averaging 1.57 on the five-point scale. Humans rated handwritten notes most acceptable at 4.36, while agents assigned that distinction to AI-generated notes at 3.75. Critically, agents ranked handwritten notes near the bottom of the scale at 2.45, placing them below getting interview questions beforehand, which is an ordering humans reversed entirely. Human ratings thus followed a meaningfully different ordinal pattern from agents, with greater dispersion across all six behaviors, reflecting genuine individual disagreement about where the line between legitimate optimization and dishonest conduct falls. This distinction between preparation-stage AI use and live substitution maps onto the pattern observed in Question 1, where AI practice was seen as common and AI answer generation as rare; taken together, the two patterns suggest that human respondents hold a coherent, if largely implicit, theory of what separates acceptable optimization from dishonest conduct, one that centers on the substitution of real-time judgment rather than AI assistance as such. Agents do not share that ordering, a discrepancy examined further in Section 4.5. Figure 6 illustrates the acceptability rating distributions by behavior.




Figure 6: Participants rated the acceptability of six interview-related behaviors on a scale from 1 (completely unacceptable) to 5 (completely acceptable). The agent means are plotted as red diamonds for reference, while the human ratings are shown as boxplots to capture their distributional characteristics. Note that “Unsure” responses are not included in the analysis. “Live Assist Behind the Scenes” received the lowest ratings ($\bar{x} = 1.573$, $\tilde{x} = 1.00$, $s = 0.901$), with little variability, suggesting near-unanimous disapproval. “Handwritten Notes” received the highest ratings ($\bar{x} = 4.361$, $\tilde{x} = 5.00$, $s = 1.068$), indicating broad acceptability. “AI-Generated Notes” fell just above the scale midpoint ($\bar{x} = 3.540$, $\tilde{x} = 4.00$, $s = 1.467$), though with notable variability suggesting divided opinion. “Questions Beforehand” received moderate ratings ($\bar{x} = 3.130$, $\tilde{x} = 3.00$, $s = 1.40$), while “Search Engine During Interview” ($\bar{x} = 2.76$, $\tilde{x} = 2.00$, $s = 1.282$) and “AI During Interview” ($\bar{x} = 2.193$, $\tilde{x} = 2.00$, $s = 1.183$) were rated as generally unacceptable. Overall, participants drew a clear distinction between preparation-based behaviors, which were viewed more favorably, and active in-interview assistance, which was largely disapproved of.

Among respondents who rated the process as “Somewhat Unfair” or “Extremely Unfair,” acceptability scores were meaningfully higher for getting interview questions beforehand than among those who rated the process as fair, as skeptic respondents rated it higher at 3.52 compared to 2.70 for those who rated the process as fair or neutral. This difference is substantively notable because it shifts the behavior’s mean rating from the unacceptable range into the acceptable range. The gap was smallest for AI assistance during the interview, which remained relatively unacceptable even among the most skeptical respondents. This suggests that moral disengagement operates at the margins of ethical judgment rather than collapsing it wholesale: unfairness perceptions expand the range of behaviors candidates consider acceptable optimization without dissolving resistance to its most unambiguous forms. The most extreme cheating behaviors appear to retain their illegitimacy even when candidates feel the system has treated them poorly. Figure 7 illustrates this cross-tabulation.




Figure 7: Cross-tabulation of fairness ratings (Q21) against acceptability scores for getting interview questions beforehand (Q3a). Respondents who rated the process as “Somewhat Unfair” or “Extremely Unfair” assigned higher acceptability scores to this behavior than those who rated it as “Somewhat Fair” or “Extremely Fair,” with mean scores of 3.52 and 2.70 respectively. This pattern supports the moral disengagement hypothesis that perceived unfairness can license greater acceptance of behaviors that might otherwise be viewed as unethical, though the effect appears to be behavior-specific rather than general.

A Spearman correlation between perceived fairness (Q21) and willingness to cheat when a peer scores high (Q13a) revealed a statistically significant negative relationship ($r = -0.175$, $p = 0.041$), consistent with the directional prediction of H2. Inspection of group means supported the directional pattern among adequately sized groups: participants who perceived the situation as extremely unfair reported the highest mean willingness ($\bar{x} = 3.48$), followed by those who perceived it as somewhat unfair ($\bar{x} = 3.24$), neither fair nor unfair ($\bar{x} = 3.24$), and somewhat fair ($\bar{x} = 2.94$). The extremely fair group ($n = 2$) was excluded from interpretation due to insufficient sample size. The correlation clears the conventional significance threshold, though the effect size is modest, and the cross-sectional design precludes causal inference. Future research with larger and more evenly distributed samples would

allow for a more robust test of this hypothesis. Figure 8 plots this relationship.




Figure 8: Fairness perception and willingness to cheat among human respondents, shown as a scatterplot. An OLS regression line is included to illustrate the overall negative trend. Each point represents an individual respondent's fairness rating (x-axis) and their willingness to cheat under high peer cheating conditions (y-axis). The plot shows a statistically significant downward slope ($r = -0.175$, $p = 0.041$), indicating that lower fairness perceptions are associated with higher willingness to cheat, consistent with moral disengagement theory. The dispersion of points and the presence of outliers suggest that other factors also influence willingness.

The open-ended responses to Question 6 provide qualitative evidence about the reasoning behind cheating that the Likert items cannot capture. Among the 73 responses received, 18% (n=13) invoked themes of competitive pressure or desperation, describing cheating as a rational response to a job market candidates perceive as stacked against them. Several respondents framed AI use explicitly in terms of fairness restoration, arguing that if employers use AI to screen candidates, candidates are entitled to use AI to respond. One respondent put it plainly: the potential reward for cheating far outweighs the cost of a rejection that was already likely. Another drew a direct equivalence with elite hiring networks, arguing that AI assistance is no different from leveraging personal connections. This framing maps directly onto Bandura (1999)’s moral disengagement mechanism: rule-breaking is licensed by the perception that the rules themselves are applied unfairly. A further 8% (n=6) of substantive responses discussed AI use as a cultural shift rather than a moral failing, with one respondent citing ”cultural norms surrounding general AI use in academic and professional settings” and others describing increasing dependence on AI as an environmental condition rather than an individual choice. Roughly 4% (n=3) explicitly cited low detection risk as a motivating factor, consistent with the deterrence framing in H3, though the small count warrants caution. Agent responses to the same open-ended item clustered around a narrower set of justifications, predominantly desperation and competition, with little of the fairness-restoration framing or cultural norm language that appeared in human responses. The divergence mirrors the pattern observed in the quantitative items: agents approximate the surface-level distribution of motivations but do not replicate the more elaborated moral reframings through which human candidates reconstruct conduct they might otherwise consider unacceptable.

The evidence supports H2. Perceived unfairness was associated with higher acceptability of getting interview questions beforehand, shifting mean ratings from unacceptable among fair-process respondents to acceptable among skeptical ones, and the directional pattern between fairness perception and willingness to cheat held across adequately sized groups. The Spearman correlation between fairness and willingness reached conventional statistical significance ($r = -0.175$, $p = 0.041$). Importantly, the effect was concentrated at the margins of ethical judgment: resistance to the most unambiguous forms of cheating persisted even among the most skeptical respondents, a pattern consistent with Bandura (1999)’s account of moral disengagement as operating selectively rather than wholesale. The correlational design cannot establish causation, and the modest effect size warrants caution in extrapolating the magnitude of the relationship. The behavior-specificity of the effect is itself a finding that the original hypothesis does not anticipate and that future work should examine directly.

4.3 H3: Deterrence

H3 predicts that lower confidence in companies’ ability to detect and act on cheating will be associated with greater willingness to engage in it. When candidates believe that the expected cost of cheating is low, whether because detection is unlikely or because consequences are unlikely to follow even when cheating is suspected, the deterrence mechanism that might otherwise restrain behavior is weakened. Four survey items bear on this hypothesis: Question 23, which asked how confident respondents were that companies can effectively detect cheating in virtual interviews; Question 24, which asked how confident they were that companies act on suspected cheating; Question 13a, which

measured willingness to use unauthorized assistance under high peer cheating conditions; and Question 25, which elicited views on the appropriate consequence for AI-assisted cheating.

Both samples converged on low confidence in detection and enforcement, though the human distribution skewed more sharply toward the floor. Among human respondents, 44.5% (n=61) reported being “Not Confident at All” in companies’ ability to detect cheating in virtual interviews, and 35.8% (n=49) were “Slightly Confident,” leaving only 19.7% above that threshold. Confidence in enforcement was somewhat higher but followed the same general pattern, with 30.7% (n=42) reporting “Not Confident at All” and 32.8% (n=45) “Slightly Confident.” Agents clustered in the middle range on both items, with the plurality selecting “Moderately Confident” on detection and splitting between “Slightly” and “Moderately Confident” on enforcement; notably, no agent selected “Not Confident at All” on either item. Taken together, the data suggest that both humans and agents regard current detection and enforcement mechanisms as inadequate, but humans are more willing to express outright skepticism, while agents tend to assign at least partial credit to institutional capacity. Figure 9 illustrates the distribution of confidence ratings for both samples.




Figure 9: Distribution of confidence ratings for companies’ detection and enforcement capabilities among human respondents and agent respondents. Both samples show a general lack of confidence, with human respondents skewing more heavily toward “Not Confident at All” and agents clustering around “Slightly Confident” and “Moderately Confident.” This suggests that while both groups share a perception of institutional inadequacy, humans are more likely to express outright skepticism, whereas agents tend to assign at least some degree of confidence to current mechanisms.

Cross-tabulating detection confidence against Question 13a responses reveals no consistent gradient between perceived detection risk and stated willingness to cheat among human respondents. Willingness rates across confidence levels were as follows: 50.8% among those “Not Confident at All” ($n=61$), 63.3% among those “Slightly Confident” ($n=49$), 50.0% among those “Moderately Confident” ($n=24$), and 50.0% among those “Very Confident” ($n=2$). Rather than rising monotonically as detection confidence falls, willingness peaks among the “Slightly Confident” group and then returns to roughly the same level across the remaining categories. A Spearman correlation between detection confidence and willingness confirmed the absence of a systematic relationship ($r = 0.054$, $p = 0.528$), which Figure 10 visualizes. Under the no-peer condition (Q14a), willingness rates were uniformly low across all confidence categories, ranging from 6.6% among those “Not Confident at All” to 4.2% among those “Moderately Confident” ($r = -0.005$, $p = 0.956$), suggesting that when social context is absent, detection confidence has no purchase on willingness regardless of its direction. These results do not support a straightforward deterrence account, in which lower perceived detection risk would be expected to produce higher stated willingness. The relationship between what respondents believe companies can detect and what they say they would do appears, at least in this sample, to be more weakly coupled than deterrence theory would predict. Whether this reflects genuine indifference to detection risk, social desirability suppression of true willingness at higher confidence levels, or simply insufficient statistical power given the small cell sizes at the upper end of the confidence scale remains an open question.




Figure 10: Cross-tabulation of detection confidence (Q23) against willingness to cheat under high peer cheating conditions (Q13a) among human respondents. The plot shows the percentage of respondents reporting “Somewhat Likely” or “Extremely Likely” to cheat within each confidence category. Contrary to deterrence theory predictions, willingness does not rise monotonically as detection confidence falls; instead, it peaks among the “Slightly Confident” group and returns to similar levels across the remaining categories. A Spearman correlation confirmed the absence of a systematic relationship ($r = 0.054$, $p = 0.528$).

Question 25, which asked about appropriate consequences for AI-assisted cheating, adds a second dimension to the deterrence picture: not just how candidates perceive detection risk, but how seriously they believe cheating should be treated. Among agents, all 150 selected automatic disqualification from the current application as the appropriate consequence for AI-assisted cheating, a response that is both decisive and uniform. Human respondents showed considerably more variation. While automatic disqualification remained the modal response ($n=101$), 11.70% ($n=16$) called for a permanent ban from applying to the company and 3.60% for an industry-wide ban. At the same time, 11.0% ($n=15$) called for lighter consequences including no penalty, on the grounds that AI use is a legitimate adaptation to modern hiring. The agent uniformity on this item is analytically significant. All 150 agents converging on the same moderate response suggests that these models resolve genuine normative disagreement into a single consensus position. The real candidate population does not share that consensus, and a tool that predicts it would misrepresent the distribution in ways directly relevant to enforcement policy design. When cross-tabulated with Question 13a, respondents who favored lighter consequences in Question 25 were 25% more likely to report willingness to cheat than those who favored disqualification or bans, suggesting that consequence severity preferences and behavioral willingness are correlated in the direction deterrence theory predicts. Figure 11 illustrates the distribution of responses to Question 25 and their relationship with willingness to cheat.

The evidence does not support H3. Contrary to the prediction that lower detection confidence would produce greater willingness to cheat, no consistent gradient emerged from the cross-tabulation of detection confidence against Question 13a responses, and the Spearman correlation confirmed the absence of a systematic relationship ($r = 0.054$, $p = 0.528$). Confidence in enforcement showed the same flat pattern. What the data do support is a narrower version of the deterrence logic: respondents who favored lighter consequences for AI-assisted cheating were meaningfully more likely to report willingness to cheat than those who favored disqualification or bans. This suggests that deterrence, to the extent it operates in this context, runs through perceived penalty magnitude rather than perceived detection probability. The current design cannot establish which component matters more, but the distinction is one that future research, and enforcement policy, would do well to take seriously.

4.4 H4: Experience and Normalization

H4 predicts that greater prior interview experience will be associated with higher normalization of dishonest behavior. Repeated exposure to competitive hiring may gradually erode categorical ethical commitments, replacing them with more contextual or instrumental reasoning. Two theoretical accounts predict this pattern. Moral disengagement theory holds that repeated exposure to morally ambiguous situations weakens ethical self-regulation over time, as individuals develop cognitive reframings that allow them to suspend judgment they would otherwise apply (Bandura, 1999). Social learning accounts offer a related but distinct mechanism: behavioral norms are updated through observation of what peers actually do, and more experienced candidates have had more opportunities to observe or hear about cheating. Both accounts predict an experience gradient in normalization, but they differ in where the effect should appear: moral disengagement predicts shifts in acceptability ratings, while social learning predicts shifts in peer cheating estimates. Four survey items bear on this hypothesis: Question 19, which recorded




Figure 11: Distribution of responses to Question 25, which asked about appropriate consequences for AI-assisted cheating. Among human respondents, 73.7% (n=101) selected automatic disqualification from the current application, 11.7% (n=16) called for a permanent ban from applying to the company, 3.6% (n=5) called for an industry-wide ban, and 11.0% (n=15) called for lighter consequences including no penalty. Among agents, all 150 selected automatic disqualification. The human distribution shows considerable variation in consequence preferences, while the agent distribution is uniform, suggesting that agents may fail to capture the normative diversity present in the candidate population. Cross-tabulation with willingness to cheat revealed that respondents favoring lighter consequences were more likely to report willingness to cheat, consistent with deterrence theory predictions.

the number of professional interviews each respondent had completed; Question 13a and Question 3, which measured willingness and acceptability respectively; and Question 2, which captured peer cheating estimates. These are supplemented by behavioral profile cross-tabulations against field of study, which allow for a test of whether normalization is experience-driven, field-specific, or both.

Respondents were grouped into five experience bands: 0 to 1 interviews, 2 to 6, 7 to 10, 11-20, and 21+. These cut points were chosen to distinguish between respondents with minimal exposure to formal hiring, those with moderate experience typical of undergraduate and early graduate job-seeking, and those with substantial experience in multiple competitive cycles. The distribution of human respondents across these bands was [38, 113, 41, 32, 25] respectively. To ensure stronger cell sizes when cross-tabulating with other measures, I combine the 7-10, 11-20, and 21+ groups together, for a final distribution of [38, 113, 98]. Figure 12 illustrates the distribution of respondents across these three experience bands.

The relationship between experience and perceived peer cheating shows a clear directional gradient. Respon-




Figure 12: Distribution of human respondents across experience bands based on the number of professional interviews completed. The bands are defined as follows: 0 to 1 interviews ($n=38$), 2 to 6 interviews ($n=113$), and 7 or more interviews ($n=98$). The distribution shows a concentration in the moderate experience category (2-6 interviews) and higher experience category (7 or more interviews), with a smaller but meaningful population in the low experience category (0-1 interviews). This distribution allows for analysis of how interview experience correlates with perceptions of peer cheating, acceptability ratings, and willingness to cheat, providing insight into the potential normalization effects of competitive exposure. Note that in the survey, a more granular set of experience bands was collected, but for analytical purposes, the higher experience categories have been combined to ensure sufficient cell sizes for cross-tabulation.

dents in the 0 to 1 interview band estimated on average that 35.29% of job-seekers they know have cheated, while respondents in the 2 to 6 band estimated 34.65%, and those with 7 or more interviews estimated 39.71%. This upward gradient is consistent with a social learning account: more experienced candidates have larger reference networks, and within those networks the probability of having encountered or heard about a cheating incident increases with exposure. Whether the gradient reflects genuinely higher cheating rates in the networks of experienced candidates or instead reflects the accumulation of salient cheating incidents that progressively revise norm perceptions upward is a question this data cannot settle. The distinction matters because the first account implies that experienced candidates simply have more accurate information, while the second implies a systematic upward bias in norm perception that grows with exposure.

Cross-tabulating experience band against Question 13a willingness provided no support for H4. Indeed, the direction of the finding runs contrary to the prediction: the least experienced group reported the highest willingness (60.5%, $n = 38$), while the most experienced group fell in between (56.1%, $n = 98$), and the moderate experience group reported the lowest willingness (52.2%, $n = 113$). The overall association was negligible and non-significant (Spearman $r = 0.028$, $p = 0.655$). Controlling for perceived peer cheating estimates from Question 2 did not alter this picture (partial $r = 0.013$, $p = 0.838$). Rather than accumulating tolerance through repeated competitive

exposure, inexperienced candidates appear at least as susceptible to peer norm pressure as those with more hiring cycles behind them. One interpretation consistent with this pattern is that the ethical disillusionment process, if it operates at all, requires a longer career horizon than this sample spans; another is that first-exposure candidates are especially responsive to peer cues precisely because they lack the reference experience needed to calibrate an independent ethical standard.

The acceptability ratings from Question 3 add texture to this picture, though again not in the direction H4 anticipated. More experienced respondents assigned *lower* acceptability scores to five of the six behaviors than the least experienced respondents, with the largest declines for using AI during the interview ($\Delta = -0.36$) and using a search engine during the interview ($\Delta = -0.35$). The sole exception was AI-generated notes, which experienced respondents rated slightly more acceptable ($\Delta = +0.10$). Rather than becoming more permissive as competitive exposure accumulates, experienced candidates appear, if anything, marginally more conservative across most forms of unauthorized assistance. These differences are small in magnitude, however, and should be read as descriptive; absent significance testing they are consistent with little or no genuine experience effect, echoing the null association observed for behavioral willingness in Question 13a.

Experience alone does not fully account for these patterns. To examine whether normalization reflects competitive culture at the field level rather than individual exposure, respondents were categorized into willingness profiles based on their responses to Questions 13a and 14a: Unwilling, meaning unlikely under both conditions ($n = 71$, 28.5%); Conditionally Willing, meaning unlikely under no-cheating conditions but somewhat likely otherwise ($n = 107$, 43.0%); Indifferent, meaning neither likely nor unlikely regardless of conditions ($n = 8$); Generally Willing, meaning likely under either condition ($n = 14$, 5.6%); and a Contrarian profile (likely under no-cheating conditions but unlikely otherwise), which had no members. A residual Mixed category ($n = 49$, 19.7%) captured response patterns not fitting these definitions. Figure 13 visualizes the distribution of human respondents across these willingness profiles. Figure 14 plots experience band composition within each profile.




Figure 13: Distribution of human respondents across willingness profiles based on their responses to Questions 13a and 14a. The profiles are defined as follows: Unwilling (unlikely under both conditions, $n=71$), Conditionally Willing (unlikely under no-cheating conditions but somewhat likely otherwise, $n=107$), Indifferent (neither likely nor unlikely regardless of conditions, $n=8$), Generally Willing (likely under either condition, $n=14$), and a Contrarian profile (likely under no-cheating conditions but unlikely otherwise, $n=0$). A residual Mixed category ($n=49$) captures response patterns not fitting these definitions. This distribution allows for analysis of how willingness to cheat varies across different social contexts and how it correlates with field of study and other demographic variables, providing insight into the potential normalization effects of competitive culture at the field level.




Figure 14: Experience distribution across willingness profiles. Each profile contains a mix of respondents across the experience bands (0-1 interviews, 2-6 interviews, 7 or more interviews), suggesting that experience does not appear to be a strong predictor of profile membership. This pattern indicates that willingness to cheat is not simply a function of accumulated interview experience, but may instead be influenced by other factors such as field-specific competitive culture or individual differences in ethical reasoning. The lack of a clear experience gradient across profiles challenges the notion that normalization of cheating behavior is primarily driven by repeated exposure to competitive hiring situations.

Software Engineering, AI, and Data Science respondents were notably overrepresented in the willing profiles relative to their share of the sample, comprising 28 of 107 Conditionally Willing respondents and 6 of 14 Generally Willing respondents. Education respondents, by contrast, skewed toward the Unwilling profile (two-thirds Unwilling). Other fields were more mixed: Healthcare and Finance respondents, despite their visible presence in the willing groups in absolute terms, were in fact majority Conditionally Willing within their own ranks (60% and 62% respectively), reflecting sample size as much as disposition. Figure 15 illustrates the mix of industries within each profile. Academic level showed no predictive power ($\chi^2 = 13.850$, $df = 21$, $p = 0.876$), with respondents distributed roughly proportionally across profiles regardless of year of study, suggesting that whatever field-level variation exists is not a function of years of academic training.




Figure 15: Left: distribution of willingness profiles within each industry, showing the proportion of respondents in each profile category (Unwilling, Conditionally Willing, Indifferent, Generally Willing, Mixed) for each field of study. Right: Distribution of industries within each profile. Software Engineering, AI, and Data Science respondents were notably overrepresented in the willing profiles relative to their share of the sample, while Education respondents skewed toward the Unwilling profile. Healthcare and Finance respondents were majority Conditionally Willing within their own ranks.

The Generally Willing group warrants particular attention. While small in number, these respondents reported willingness to use unauthorized assistance regardless of what peers were doing, making peer behavior irrelevant as a moral anchor. They were distinguished most clearly not by deterrence perceptions but by their beliefs about prevalence and acceptability: they estimated peer cheating at 59% on average (Question 2), far above the Conditionally Willing (37%) and Unwilling (29%) groups, and rated the behaviors as more acceptable overall (Question 3 mean 3.57 versus 3.04 and 2.60). Their detection-confidence scores (Question 23) were marginally the lowest of the three groups, though their enforcement-confidence scores (Question 24) were not, so the data do not support a clean “low perceived consequences” account. The more consistent interpretation is that this group’s willingness tracks a belief that cheating is both widespread and broadly acceptable, rather than a calculated read of low enforcement risk. Distinguishing this group from the Conditionally Willing, whose willingness is contingent on peer behavior, has practical implications for how enforcement and norm-based interventions might be targeted. Figure 16 compares key measures across profiles.




Figure 16: Comparison of key measures across the Unwilling, Conditionally Willing, and Generally Willing profiles. The Generally Willing group estimated peer cheating at 59% on average, far above the Conditionally Willing (37%) and Unwilling (29%) groups, and rated the behaviors as more acceptable overall (Question 3 mean 3.57 versus 3.04 and 2.60). Their detection-confidence scores were marginally the lowest of the three groups, though their enforcement-confidence scores were not, suggesting that their willingness is more closely tied to beliefs about prevalence and acceptability than to perceptions of enforcement risk. Note that the peer cheating estimates are scaled down by a factor of 10, so the visual comparison should be interpreted with that in mind.

The evidence supports the view that competitive culture at the field level, rather than individual experience, is the more salient driver of normalization. The overrepresentation of certain fields in the willing profiles suggests that industry-specific norms and pressures may shape candidates' ethical reasoning in ways that are not simply a function of how many interviews they have been through. The lack of a clear experience gradient across profiles challenges the notion that normalization is primarily driven by repeated exposure to competitive hiring situations, and instead points to the importance of field-specific cultures and norms in shaping candidates' perceptions and behaviors around cheating.

4.5 H5: Human-AI Divergence

H5 predicts that AI agents will systematically underestimate willingness to cheat and overestimate ethical uniformity relative to human respondents. This section treats the divergence itself as the primary object of analysis, asking




Figure 17: Mean divergence and variance ratio across all H5-relevant items, grouped by social contingency type. The left panel plots the gap between human and agent means (human minus agent), such that positive values indicate human respondents scored higher; the right panel plots the ratio of human to agent standard deviations, where values above one indicate greater dispersion among human respondents. Items are color-coded by the type of social judgment involved: red for items driven primarily by social context and competitive pressure, orange for the acceptability hierarchy items from Q3, green for items showing close convergence between samples, and purple for deterrence-related confidence items. The dashed reference line in the right panel marks equal dispersion between samples. Detailed discussion of each item appears in the sections that follow.

not simply whether agents and humans differ but where the divergence is sharpest, where it narrows, and what distinguishes the two. Figure 17 organizes the evidence bearing on that prediction across all items in the study. The left panel plots mean divergence by item, with positive values indicating that human respondents scored higher than agents; the right panel plots the ratio of human to agent standard deviations, where values above one indicate that humans were more dispersed. Both panels sort items by the magnitude of divergence and color-code them by the type of judgment involved. The figure makes visible a pattern that the sections below develop in detail: divergence is not uniform across items but structured by the degree to which a given judgment depends on social context, competitive exposure, and individual variation.

The most substantial divergences cluster on items where social context is the primary driver of response. As documented in H1, agent estimates of peer cheating prevalence (mean 12.7%, clustered almost entirely at 5, 10, and 15%) differed from human estimates not merely in magnitude but in the reasoning process they reflect: agents produce something like an abstract prior, while human respondents draw on actual social networks shaped by their competitive environments. This difference, the largest in the study at a gap of 24.0 percentage points on Question 2, is not incidental. It is the mechanism through which all downstream divergences propagate. Agents who underestimate peer cheating will necessarily underestimate the normative pressure that peer cheating exerts,

and that underestimation compounds across every item that asks how candidates respond to competitive context.

The scenario-based items make this propagation visible. The within-respondent shift from Question 14a to Question 13a was meaningfully positive for human respondents (mean shift of 1.606 scale points, from 1.622 to 3.229) and substantially attenuated for agents (mean shift of 0.933), producing a context-sensitivity gap of 0.673. This insensitivity to situational framing is the clearest behavioral signature of what H5 identifies. Human respondents update their stated willingness in response to perceived competitive pressure; agents apply a stable ethical judgment regardless of the context they are placed in. Agents are not simply more ethical than humans in some absolute sense. They are less sensitive to the social inputs that make ethical judgment context-dependent. That context-dependence is precisely what descriptive norm theory, moral disengagement theory, and deterrence theory all require.

The pattern is more nuanced on items where social contingency is lower. The ordinal acceptability hierarchy documented in H2, where live assistance rated least acceptable and handwritten notes most acceptable among human respondents, held across both samples, with a rank correlation of $r = 0.886$ ($p = 0.019$). What agents cannot replicate is the dispersion around that hierarchy: human acceptability ratings showed a mean standard deviation of 1.216 across behaviors, compared to an agent SD of approximately 0.150, a ratio of 8.1 to 1. This is a consistent pattern. Wherever a stable, socially shared norm exists, agents approximate it reasonably well. Wherever outcomes are driven by individual variation in competitive experience, social exposure, or situational framing, agent predictions systematically overestimate uniformity.

The fairness items from Question 21 represent one of the closer alignments between the two samples, as reported in H2. Human respondents averaged 2.496 and agents 3.213 on the fairness scale, a gap of -0.717 , driven by agents clustering more tightly around moderate responses. Neither group selected extreme options in meaningful numbers. This convergence is interpretively useful: fairness assessments of an institutional process may be less socially contingent than willingness judgments, and agents assigned to profiles with varied attitudinal characteristics appear better equipped to replicate considered, moderate evaluations than to simulate reactive, pressure-sensitive reasoning. The convergence on Question 21 is not evidence that agents perform well in general; it is evidence that their performance is conditioned on the type of judgment required.

Question 22, which asked how likely respondents were to report a peer who had cheated, produced the most striking convergence of the study. Both human respondents and agents reported being overwhelmingly unlikely to report cheating peers: 82.5% of humans and 98% of agents selected “Extremely Unlikely” or “Somewhat Unlikely,” for a gap of just 0.100 scale points. This convergence is notable precisely because it is not predicted by H5’s general expectation of agent optimism. Agents did not predict that candidates would be eager to report peers; they correctly anticipated a strong norm of peer loyalty that holds independent of personal ethical assessments of cheating. This suggests that some social norms, particularly those governing ingroup solidarity and perceived betrayal, are sufficiently stable and widely understood that current large language models can replicate them without the kind of real social pressure that produces divergence elsewhere. The convergence on Question 22 strengthens confidence that the divergences observed elsewhere are not artifacts of general agent conservatism but reflect genuine failures

to simulate context-sensitive reasoning.

Both samples agreed that institutional confidence was low: no human or agent respondent selected “Very Confident” or “Extremely Confident” on either item. The divergence lies in degree rather than direction. On detection (Q23), 44.5% of human respondents selected “Not Confident at All,” a response that no agent chose on either item; on enforcement (Q24), 30.7% of human respondents selected the same floor response, again with no agent equivalent. The detection gap (-0.832 on a five-point scale) was the larger of the two, with the enforcement gap a more modest -0.282. Agents capture the general direction of candidate skepticism about employer capabilities but systematically underestimate its depth, compressing distributions that extend to the floor of the confidence scale into a narrower middle range. On Question 25, as reported under H3, all 150 agents converged uniformly on automatic disqualification while human respondents distributed across a range from no penalty to industry-wide bans. This uniformity is a further instance of the ethical consensus overestimation that H5 predicts.

The cumulative picture across these items supports H5. Agents systematically underestimate willingness to cheat, overestimate ethical uniformity, and fail to replicate the contextual sensitivity that makes descriptive norm effects, moral disengagement, and deterrence relevant as behavioral mechanisms. The areas of convergence, reporting norms and the ordinal hierarchy of cheating acceptability, are instructive: agents can replicate stable, socially consensual norms that do not require situational updating, but they cannot simulate the reactive, pressure-sensitive reasoning that produces the most behaviorally consequential variation in human responses. The divergence is not uniform; it is structured by the degree to which a given judgment depends on social context, competitive exposure, and individual variation in experiential history. That structure is itself a finding, and it points toward a more precise characterization of current model limitations than a simple “agents are more ethical” account would provide.

5 Shortcomings and Limitations

This study carries limitations in two areas: those specific to the agent-based component and those arising from the human survey design. Each is addressed in turn.

5.1 Agentic Survey Limitations

To begin with the agent-based survey, a fundamental constraint concerns the behavioral consistency of large language models in their current state. Agents tend to produce highly uniform responses across administrations, which diverges from the variability characteristic of real human survey data. Human respondents are subject to fatigue, retention loss, and mood variation; these sources of noise are not incidental but are constitutive of how real attitudes are expressed. They may also hold internally inconsistent views, rating a behavior “extremely unacceptable” in one item while describing a comparable behavior as “acceptable” in another. Their responses are further shaped by contextual factors including physical setting, personal ethical commitments, and prior experience with AI tools or

job searching. Current AI agents cannot replicate these psychological dynamics without being explicitly instructed to do so, and such instructions introduce their own distortions by pushing agents away from naturalistic response patterns. The H5 analysis documents this empirically: across social-context-dependent items, human standard deviations exceeded agent standard deviations by a ratio of roughly 8 to 1, and agents showed no meaningful within-respondent shift between the high-peer and no-peer cheating conditions. The limitation described here is therefore not merely a methodological caveat offered in anticipation of problems; it is a pattern the data confirm.

A related limitation concerns prompt sensitivity. Agent behavior is highly dependent on how prompts are constructed: small differences in framing or persona instructions can shift response distributions substantially, meaning that findings may reflect the researcher’s prompting choices as much as the underlying constructs of interest. In this study, prompts were designed to be deliberately minimal, avoiding strong persona cues or attitudinal anchors, in order to reduce researcher-induced bias. This approach partially addresses the consistency problem but does not eliminate it, and future work using agent-based methods would benefit from systematic prompt variation as a robustness check.

Taken together, these constraints mean that agent-generated responses should be interpreted as idealized or theoretical approximations of human opinion rather than as empirical estimates of real-world attitudes. The agent data are best treated as predictive in nature, useful for identifying directional patterns and generating hypotheses, but not for extrapolating population-level trends.

5.2 Human Survey Limitations

The human survey carries its own set of limitations, which fall into three broad categories: sample representativeness, self-report bias, and design constraints.

The most significant concern is the composition of the sample. Respondents were drawn from a primarily college student population, which introduces several sources of non-representativeness. Students tend to be younger, more formally educated, and more familiar with technology than the general workforce, meaning their attitudes toward AI in hiring may be more favorable, or at least more informed, than those of older or less tech-literate job seekers. Equally important, many student respondents have limited direct experience with the job search process, meaning their survey responses reflect anticipated rather than experienced behavior. The H4 analysis bears this out: roughly 62% of respondents fell within the 0 to 6 interview experience band, indicating that the sample skews toward candidates projecting forward rather than drawing on accumulated competitive exposure. Findings should therefore be interpreted as reflecting the attitudes of a relatively tech-savvy, early-career population rather than job seekers broadly. Future research would benefit from sampling across age groups and experience levels, particularly respondents who have recently undergone AI-screened application processes.

A second limitation concerns self-report bias. Attitudes toward algorithmic hiring are a socially charged topic, and respondents may have adjusted their answers, consciously or not, to align with what they perceived as the

ethically correct position. This social desirability effect could lead to an overstatement of concern about AI fairness and an understatement of practical acceptance. Relatedly, expressed attitudes in a survey context do not necessarily predict behavior: a respondent who characterizes AI resume screening as unfair may nonetheless apply through AI-screened platforms without hesitation. This attitude-behavior gap is well-documented in consumer ethics research and is equally relevant in employment contexts. It cannot be fully resolved through survey design alone, though future work incorporating behavioral measures or vignette-based choice tasks would help address it.

A third limitation concerns the anonymous coding protocol used in Block 3. Respondents were assigned randomized symbols to represent yes and no responses, a design intended to lower the perceived personal stakes of disclosing sensitive peer conduct by introducing plausible deniability. In practice, however, the encoding introduced an additional layer of cognitive demand: respondents had to learn and retain their assigned symbol pair mid-survey before applying it across six consecutive items. Six respondents failed one of the two factual calibration checks embedded in this block, while passing the other. Because failing both checks would more strongly indicate disengagement, and because the symbol system’s complexity makes a single encoding error plausible even among attentive respondents, these six were retained in the main analysis. Their Block 3 responses, however, should be interpreted with some caution. Future implementations of this protocol would benefit from more prominent reminder cues or a simpler encoding scheme to reduce the risk of misapplication.

Beyond these specific concerns, several design-level constraints are worth acknowledging. The cross-sectional nature of the survey means that responses reflect attitudes at a single point in time, during a period when public awareness of AI in hiring is shifting rapidly. Opinions captured today may look meaningfully different from those collected a year earlier or later, which limits the durability of the findings. Sample size also constrains the depth of subgroup analysis: while directional differences across groups such as gender, field of study, or prior job search experience may be visible in the data, the statistical power to draw confident conclusions about these subgroups is limited. These constraints do not invalidate the findings, but they do define the scope of what can be claimed. The results speak most confidently to the attitudes of a young, tech-familiar, early-career population at a particular moment in time, and should be interpreted accordingly.

6 Broader Implications of the Results and Future Work

6.1 Theoretical Implications

The positive association between perceived peer cheating and individual willingness (H1) provides direct empirical support for focus theory’s account of descriptive norms as behavioral drivers in competitive, ambiguous contexts. The relationship holds not only at the aggregate level but in within-respondent comparisons, where stated willingness shifted sharply between the high-peer and no-peer conditions. This within-subject sensitivity weighs against a simpler account in which cheating willingness reflects fixed individual disposition; the responsiveness to situational

framing is itself the norm mechanism at work.

The support for H2 extends moral disengagement theory in a specific way: unfairness perceptions did not collapse ethical judgment wholesale but expanded it selectively, raising acceptability for borderline behaviors (getting questions beforehand) while leaving resistance to the most unambiguous violations largely intact. This behavior-specificity is not predicted by Bandura (1999)’s original account and suggests that moral disengagement in institutional settings operates at the margins rather than as a threshold shift. The implication is not that Bandura’s account is wrong but that it requires a more fine-grained specification: the relevant question is not whether unfairness licenses disengagement but which behaviors fall within the disengagement zone. Future theoretical work should examine what properties of a behavior, including its observability, its proximity to live performance, and its reversibility, determine that boundary.

The null finding for H3 is itself theoretically informative. Contrary to classical deterrence theory, perceived detection risk showed no systematic relationship to stated willingness: candidates who expressed the highest skepticism about institutional detection capacity were not meaningfully more willing to cheat than those who assigned it moderate credibility. What did correlate with willingness was preference for lighter consequences. This pattern suggests that deterrence in this context operates primarily through perceived penalty magnitude rather than perceived detection probability. One explanation consistent with the data is structural: when nearly all candidates perceive detection as unlikely, variance in that dimension collapses, leaving consequence severity as the only differentiating deterrence input. This is not a finding that standard deterrence models anticipate. In domains where detection is widely regarded as impossible, the standard model’s two-component structure collapses to one, and the variance needed to test the detection component disappears.

The null finding for H4 deserves explicit theoretical treatment. The ethical disillusionment concept, introduced to predict that accumulated interview experience would erode categorical ethical commitments, found no support in the data. Inexperienced candidates were at least as susceptible to peer norm pressure as those with many hiring cycles behind them, and experienced candidates were marginally more conservative across most forms of unauthorized assistance. Two interpretations are possible. One is that the disillusionment process requires a longer career horizon than this sample spans, with meaningful normalization visible only after sustained professional exposure rather than repeated undergraduate job-seeking. The other is that the relevant driver of normalization is competitive culture at the field level rather than individual experience accumulated over time, which would explain why Software Engineering and Data Science respondents skewed toward willing profiles regardless of how many interviews they had completed.

6.2 Practical Implications

As AI tools become more sophisticated, the boundary between using them for preparation and claiming their outputs as one’s own grows increasingly contested. That a majority of respondents found AI-generated notes highly

acceptable suggests employers may face growing pressure to treat AI fluency as a baseline competency rather than a disqualifying advantage, consistent with broader industry trends in which LLM proficiency is listed as a desired skill in technical job postings (LinkedIn Economic Graph, 2023). This normative trajectory has historical precedent: calculators were once prohibited in professional assessments and are now standard. If AI tools follow the same trajectory, what is currently treated as cheating may be reframed as applied tool use, shifting the design challenge from policing AI use to distinguishing candidates who use it fluently from those who rely on it entirely. The data establish both the direction of current tolerance and where it ends: preparation-stage AI use is already normalized, live substitution is not, and the distance between those two thresholds is what enforcement policy must currently defend.

Respondents' skepticism about detection and their sense of process unfairness are worth reading together rather than as separate grievances. Together they describe a candidate pool that doubts both the integrity of the process and the ability of its gatekeepers to enforce it. This combination is consequential because it may erode voluntary compliance independent of any actual increase in cheating. When candidates distrust the process and expect enforcement to be weak, even those who would otherwise abstain from misconduct may conclude that unilateral restraint puts them at a competitive disadvantage. Addressing these perceptions separately, by improving fairness without addressing enforcement credibility or vice versa, is likely to be less effective than interventions that target both.

The fairness data carry a specific implication for recruitment design. If candidates systematically overestimate how widespread AI cheating is, as Baugh et al. (2021) document in high-stakes contexts more generally, then transparency interventions that convey accurate base rates could reduce uptake by correcting that misperception without any change to enforcement mechanisms. Because penalty magnitude predicted willingness where detection confidence did not, enforcement policies built around clearly communicated and consistently applied consequences are likely to outperform those concentrated on surveillance. In domains where candidates already assume detection is nearly impossible, increasing surveillance capacity does not address the operative deterrent variable.

If the patterns observed in H5 generalize to deployed contexts, AI tools used to model or predict candidate behavior in hiring will produce systematically optimistic estimates of candidate integrity under competitive pressure. A recruiter relying on AI-generated predictions of how candidates respond when they believe peers are cheating will underestimate both the prevalence of that willingness and its sensitivity to competitive context. Similarly, AI tools used to design enforcement policies will tend to predict more uniform candidate support for moderate consequences than actually exists, potentially producing policies miscalibrated in either direction: too lenient for the most skeptical candidates, or too punitive relative to those who regard AI use as legitimate adaptation. The practical implication is not to abandon AI-generated behavioral predictions but to treat them as priors that require calibration against real candidate data, particularly on items where social pressure and competitive context are the primary drivers of response. Concretely, this means validating AI-predicted willingness distributions against human survey benchmarks before using them to inform enforcement or policy design, and treating systematic underestimation of tail responses as a known bias to correct for rather than an acceptable approximation.

6.3 Directions for Future Research

Several directions follow directly from these findings. Most consequentially, the social norm mechanism proposed here remains correlational. A controlled experiment in which participants are randomly exposed to different estimates of peer AI use would provide the causal leverage needed to establish whether descriptive norms genuinely drive willingness or merely correlate with it. Tracking the same survey items longitudinally over two to three years would determine whether the normative shift documented here represents genuine attitude change or a cohort-specific snapshot. A parallel recruiter survey would add a valuable second vantage point: if employers’ confidence in detecting AI use systematically diverges from candidates’ perception of that confidence, the gap itself becomes a target for intervention. Geographic expansion would test whether these patterns generalize beyond a U.S. sample; markets with distinct professional norms around assessment and collaborative work, East Asia in particular, would be informative. Finally, a behavioral follow-through study, one that observes actual application or interview conduct rather than stated willingness, would address what is perhaps the deepest limitation of this design: the attitude-behavior gap. Candidates who report conditional willingness to cheat may not translate that willingness into action, and the degree of slippage between stated and revealed preference is itself a theoretically and practically significant quantity.

7 Conclusion

This study set out to map how early-career candidates reason about AI-assisted cheating in hiring. What it found is a normative landscape that is more structured and more contingent than either a story of rising misconduct or a story of stable ethical resistance would capture.

Candidates hold a coherent, if largely implicit, hierarchy: preparation-stage AI use is broadly tolerated, live substitution is not, and that distinction held across both human and agent samples. What the human data add, and what agents could not replicate, is the contingency of that hierarchy. Perceived peer cheating, fairness grievances, and low confidence in institutional enforcement each shift the threshold independently, and their effects compound. The candidate who believes peers are cheating, views the process as unfair, and doubts that companies will act on suspicion occupies a very different moral environment than one who holds none of those perceptions, even if their baseline attitudes are identical. Enforcement, norm communication, and process fairness are not separate policy levers; they interact. Interventions that treat them independently will underestimate how much each depends on the others.

The agent comparison sharpens the methodological contribution. Where agents performed well was precisely where social consensus was strongest: they replicated the ordinal hierarchy of cheating acceptability and correctly anticipated the strong norm against reporting peers. Where they failed was equally specific: the sharp increase in willingness under competitive pressure, the tail of deeply skeptical candidates, and the genuine disagreement

about appropriate consequences all fell outside what agents could simulate. This is not a general indictment of AI-generated behavioral predictions. It is a precise characterization of where those predictions break down. Recruiters and hiring platforms that rely on such tools should treat the contextually sensitive, socially contingent items as the specific blind spots most likely to produce a distorted picture of real candidate populations.

What this study establishes is that the negotiation over what AI assistance in hiring means and where its limits lie is already underway among candidates. Their attitudes are a shaping force in that negotiation, not simply a downstream response to employer decisions. The most consequential open question remains causal: whether manipulating peer cheating information actually changes behavior, rather than merely correlating with it, requires a controlled experiment to establish. Whether the preparation-stage versus live-substitution distinction that candidates draw so consistently in self-report holds when actual behavior is observed is perhaps the most fundamental question this work raises. And whether these patterns generalize beyond a young, tech-familiar, early-career population is a second open question with direct practical stakes. These are not loose ends; they are the questions this study has sharpened enough to make worth answering.

References

- Bandura, A. (1999). Moral disengagement in the perpetration of inhumanities [PMID: 15661671]. *Personality and Social Psychology Review*, 3(3), 193–209. https://doi.org/10.1207/s15327957pspr0303_3
- Baugh, C. M., Foster, C. A., Johnson, B. R., & D’Lauro, C. (2021). Pluralistic ignorance as a contributing factor to concussion underreporting. *Health Education & Behavior*, 49(2), 340–346. <https://doi.org/10.1177/1090198121995732>
- Becker, G. S. (1968). Crime and punishment: An economic approach. *Journal of Political Economy*, 76(2), 169–217. Retrieved May 14, 2026, from <http://www.jstor.org/stable/1830482>
- Chen, W. X., Srinivasan, S., & Zakerinia, S. (2024, December). *Displacement or complementarity? The labor market impact of generative AI* (tech. rep. No. 25-039). Harvard Business School. https://www.hbs.edu/ris/Publication%20Files/25-039_05fbec84-1f23-459b-8410-e3cd7ab6c88a.pdf
- Cialdini, R. B., Kallgren, C. A., & Reno, R. R. (1991). A focus theory of normative conduct: A theoretical refinement and reevaluation of the role of norms in human behavior. In M. P. Zanna (Ed.). Academic Press. [https://doi.org/https://doi.org/10.1016/S0065-2601\(08\)60330-5](https://doi.org/https://doi.org/10.1016/S0065-2601(08)60330-5)
- Contractor, Z., & Reyes, G. (2025, August). *Generative ai in higher education: Evidence from an elite college* (IZA Discussion Paper No. 18055). Institute of Labor Economics (IZA). Bonn. <https://www.iza.org/publications/dp18055>
- Galdin, A., & Silbert, J. (2025, November). *Making talk cheap: Generative AI and labor market signaling* [Job market paper]. Retrieved November 4, 2025, from https://jesse-silbert.github.io/website/silbert_jmp.pdf
- Hanson, M. (2025, October). College graduation statistics. Retrieved October 22, 2025, from <https://educationdata.org/number-of-college-graduates>

- Horodyski, P. (2023). Recruiter’s perception of artificial intelligence (AI)-based tools in recruitment. *Computers in Human Behavior Reports*, 10. <https://doi.org/10.1016/j.chbr.2023.100298>
- LinkedIn Economic Graph. (2023, August). *Future of work report: AI at work* (tech. rep.). LinkedIn. <https://economicgraph.linkedin.com/content/dam/me/economicgraph/en-us/PDF/future-of-work-report-ai-august-2023.pdf>
- Raghavan, M., Barocas, S., Kleinberg, J., & Levy, K. (2020). Mitigating bias in algorithmic hiring: Evaluating claims and practices. *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency*, 469–481. <https://doi.org/10.1145/3351095.3372828>
- Stantcheva, S. (2023). How to run surveys: A guide to creating your own identifying variation and revealing the invisible. *Annual Review of Economics*, 15(Volume 15, 2023), 205–234. <https://doi.org/https://doi.org/10.1146/annurev-economics-091622-010157>
- Tourangeau, R., & Yan, T. (2007). Sensitive questions in surveys. *Psychological Bulletin*, 133(5), 859–883. <https://doi.org/10.1037/0033-2909.133.5.859>
- Zhang, L., & Yench, C. (2022). Examining perceptions towards hiring algorithms. *Technology in Society*, 68, 101848. <https://doi.org/10.1016/j.techsoc.2021.101848>

A Appendix

A.1 Questionnaire

A.1.1 Questions on Frequency and Acceptability

1. How common do you think the following activities are among students you know?

Respondents were prompted with 8 of these items:

- Cheating on a significant other
- Flying a helicopter
- Drinking coffee, tea, or other caffeinated beverages at least once per week
- Performing music at a concert
- Paying someone else to take an interview for them
- Using unauthorized assistance during a course exam (e.g., prohibited internet or AI use)
- Using AI to generate full interview answers
- Using AI to practice preparing for interviews
- Having someone else write your homework essay for them
- Using hidden assistance during a remote interview

- Going to the gym consistently
- Sky diving
- Using their smartphone in the morning
- Listening to music during daily activities (working, gym, cooking, etc.)

Very Rare; Rare Uncommon; Common; Very Common

2. Out of all job-seekers that you know, what do you believe is the percentage of students that have cheated on an interview?

A slider from 0 to 100 percent

3. How acceptable or unacceptable do you consider each of the following actions during an interview?

- Using AI tools during the interview
- Using a search engine (e.g. Google) during the interview
- Getting interview questions beforehand
- Having someone else give live assist behind the scenes
- Using handwritten notes for expected and common questions
- Generating notes using AI tools

Completely Unacceptable; Somewhat Unacceptable; Neutral; Somewhat Acceptable; Acceptable; Unsure

4. Please select the "Yellow" answer choice.

Red; Blue; Yellow; Green

This question was used as an attention check.

A.1.2 Questions on Reasoning

1. Rank these 8 choices from MOST LIKELY (1) to LEAST LIKELY (8) reasons why a job-seeking student might choose to cheat in an interview.

Desperation to find a job; High competition in the market; Lack of confidence; Unfair hiring process; Online interviews reduce risk of being caught; Cultural norms around academic or professional dishonesty; Lack of guidance on what is allowed vs. not allowed; Bandwagon because others are cheating

2. Anything else you would like to add about why people cheat in interviews?

Text box

A.1.3 Questions on Awareness

For this section, respondents were given a random symbol to represent "Yes", notated here as [HAVE_SYMBOL], and another symbol to represent "No", notated here as [HAVE_NOT_SYMBOL]. These consisted of: Odd Number, Even Number, Uppercase Letter, Lowercase Letter, Color, and Object.

The next set of questions will ask you to say whether you have or have not performed certain behaviors. If the answer is YES, select the answer choice that corresponds to a [HAVE_SYMBOL]. If the answer is NO, instead select the answer choice that corresponds to a [HAVE_NOT_SYMBOL].

1. Is Chicago located in Illinois?
1; 2; A; a; Yellow; Plane
2. Is the city of Chicago in New York?
5; 8; D; d; Red; Computer
3. Do you know someone who used Artificial Intelligence to complete a homework assignment when its usage was explicitly NOT allowed?
55; 464; T; t; Phone; Purple
4. Do you know someone who used Artificial Intelligence generate interview answers for you in advance when its usage was NOT allowed?
19; 204; R; r; Pink; Pencil
5. Do you know someone who used a search engine or asked a peer to generate interview answers live when that was NOT allowed?
301; 6422; Q; q; Green; Bottle
6. Do you know someone who used Artificial Intelligence to generate responses live during a virtual interview when its usage was NOT allowed?
1527; 228; F; f; Orange; Bowl

A.1.4 Questions on Theoretical Peer Pressure

1. If you knew a HIGH NUMBER of competing candidates were cheating for a role you want, how likely would you be to also use unauthorized assistance to compete?
Extremely unlikely; Somewhat unlikely; Neither likely nor unlikely; Somewhat likely; Extremely likely; Unsure
2. How acceptable do you believe it is to cheat in this scenario?
Completely Acceptable; Somewhat Acceptable; Neutral; Somewhat Unacceptable; Completely Unacceptable; Unsure

3. If you knew NO CANDIDATES were cheating, how likely would you be to cheat?

Extremely unlikely; Somewhat unlikely; Neither likely nor unlikely; Somewhat likely; Extremely likely; Unsure

4. How acceptable do you believe it is to cheat in this scenario?

Completely Acceptable; Somewhat Acceptable; Neutral; Somewhat Unacceptable; Completely Unacceptable; Unsure

A.1.5 Questions on Demographics

1. What is your year?

Freshman; Sophomore; Junior; Senior; Graduate Student; New Grad (0-1 YOE); Experienced Hire (2+)

2. What is your current field of study?

Computer Science; Data Science / AI / Machine Learning; Electrical or Computer Engineering; Mechanical Engineering; Civil / Environmental Engineering; Chemical Engineering; Business Administration; Finance / Accounting; Economics; Psychology; Sociology / Anthropology; Political Science; Biology; Neuroscience; Chemistry; Physics; Mathematics or Statistics; Communications / Media / Journalism; Education; Graphic Design / UX / Visual Arts; Other (please specify) - text box

3. What is your current industry, or desired industry?

Software Engineering / AI / Data Science; Information Technology / Cybersecurity; Finance: Investment Banking; Finance: Asset Management / Hedge Funds; Finance: Corporate Finance / FP&A; Management Consulting; Engineering: Electrical / Hardware; Engineering: Mechanical / Robotics; Engineering: Civil / Environmental; Healthcare: Clinical / Medical; Healthcare: Biotech / Pharma / Research; Consumer Goods / Retail / E-Commerce; Media / Journalism / Communications; Entertainment / Gaming / Creative Arts; Energy / Utilities; Government / Public Sector; Nonprofit / NGOs; Education / Teaching; Other (please specify) - text box

4. What is your current employment status?

Full-time employee; Part-time employee; Internship / co-op; Self-employed / contracting; Unemployed but actively seeking work; Unemployed and not seeking work; Full-time student (no employment); Full-time student (working part-time); Prefer not to say

5. How many professional interviews have you completed?

0-1; 2-6 ; 7-10; 11-20; 21+

6. What is the primary format of your interviews?

Online; In-person; Mixed

A.1.6 Feedback and Raffle Entry

1. If you have any feedback, please feel free to share it here.

Text-box

2. Would you like to enter into the raffle to win a prize? Your answers will still remain anonymous.

Note: Completion of the optional section of the survey will earn TWO (2) additional entries into the raffle, for a total of THREE (3) entries.

Yes, I would like to enter the raffle; No, I would NOT like to be entered into the raffle

Those that respond with "Yes, I would like to enter the raffle" are redirected to a separate survey at the end of this questionnaire to receive their contact information.

3. This is the end of the required section. If you would like to provide some additional insight, then select "Yes, continue" to continue with the optional portion of the survey. If you like to go to the end of the survey, then select "No, skip to the end of the survey".

Yes, continue; No, skip to the end of the survey

A.1.7 Optional Questionnaire

1. Overall, how fair do you believe the modern job interview process is for a qualified candidate?

Extremely unfair; Somewhat unfair; Neither fair nor unfair; Somewhat fair; Extremely fair

2. How likely are you to report a friend or peer if you knew they cheated to get a job?

Extremely unlikely; Somewhat unlikely; Neither likely nor unlikely; Somewhat likely; Extremely likely

3. How confident are you that companies can effectively detect cheating in virtual interviews?

Not confident at all; Slightly confident; Moderately confident; Very confident; Extremely confident

4. How confident are you that companies act on suspected cheating?

Not confident at all; Slightly confident; Moderately confident; Very confident; Extremely confident

5. If a company discovers a candidate used AI to cheat on a critical hiring assessment, what do you believe is the appropriate consequence?

Warning, no penalty; Automatic disqualification from the current application; Permanent ban from applying to the company; No consequence, as it's part of modern technology use; Industry-wide ban

6. If a candidate secures a job primarily because of high-quality, AI-generated answers, how confident are you that they possess the genuine skills required for the role?

Not confident at all; Slightly confident; Moderately confident; Very confident; Extremely confident

A.2 AI Agent Study Responses

A.2.1 Question 1: How common do you think the following activities are among students you know?

Table 2: Question 1.a: Cheating on a Significant Other

Frequency	<i>n</i>
Very Rare	0
Rare	150
Uncommon	0
Common	0
Very Common	0

Table 3: Question 1.b: Flying a Helicopter

Frequency	<i>n</i>
Very Rare	150
Rare	0
Uncommon	0
Common	0
Very Common	0

Table 4: Question 1.c: Drinking coffee, tea, or other caffeinated beverages at least once per week

Frequency	<i>n</i>
Very Rare	0
Rare	0
Uncommon	0
Common	0
Very Common	150

Table 5: Question 1.d: Performing music at a concert

Frequency	<i>n</i>
Very Rare	0
Rare	150
Uncommon	0
Common	0
Very Common	0

Table 6: Question 1.e: Paying someone else to take an interview for them

Frequency	<i>n</i>
Very Rare	143
Rare	7
Uncommon	0
Common	0
Very Common	0

Table 7: Question 1.f: Using unauthorized assistance during a course exam (e.g., prohibited internet or AI use)

Frequency	<i>n</i>
Very Rare	0
Rare	0
Uncommon	150
Common	0
Very Common	0

Table 8: Question 1.g: Using AI to generate full interview answers

Frequency	<i>n</i>
Very Rare	0
Rare	65
Uncommon	82
Common	3
Very Common	0

Table 9: Question 1.h: Using AI to practice preparing for interviews

Frequency	<i>n</i>
Very Rare	0
Rare	0
Uncommon	10
Common	136
Very Common	4

Table 10: Question 1.i: Having someone else write your homework essay for them

Frequency	<i>n</i>
Very Rare	0
Rare	150
Uncommon	0
Common	0
Very Common	0

Table 11: Question 1.j: Using hidden assistance during a remote interview

Frequency	<i>n</i>
Very Rare	4
Rare	146
Uncommon	0
Common	0
Very Common	0

Table 12: Question 1.k: Going to the gym consistently

Frequency	<i>n</i>
Very Rare	0
Rare	0
Uncommon	5
Common	145
Very Common	0

Table 13: Question 1.l: Using their smartphone in the morning


Frequency	<i>n</i>
Very Rare	0
Rare	0
Uncommon	0
Common	0
Very Common	150

Table 14: Question 1.m: Listening to music during daily activities (working, gym, cooking, etc.)

Frequency	<i>n</i>
Very Rare	0
Rare	0
Uncommon	0
Common	0
Very Common	150

A.2.2 Question 2: Out of all job-seekers that you know, what do you believe is the percentage of students that have cheated on an interview?

Out of all job-seekers that you know, what do you believe is the percentage of students that have cheated on an interview?



Agents were asked to answer a whole number between 0 and 100 (inclusive).

A.2.3 Question 3: How acceptable or unacceptable do you consider each of the following actions during an interview?

Table 15: Mean Acceptability Ratings for Interview Cheating Behaviors

Behavior	Mean Score
Having someone else give live assist behind the scenes	1.030303
Using handwritten notes for expected and common questions	2.454545
Getting interview questions beforehand	2.548148
Using AI tools during the interview	2.780
Using a search engine (e.g. Google) during the interview	2.826
Generating notes using AI tools	3.746

A.2.4 Question 4: Please select the "Yellow" answer choice.

Table 16: Distribution of Responses for Question 4

Answer Choice	<i>n</i>
Red	0
Blue	0
Yellow	150
Green	0

A.2.5 Question 5: Rank these 8 choices from MOST LIKELY (1) to LEAST LIKELY (8) reasons why a job-seeking student might choose to cheat in an interview.

Table 17: Average Ranking for Common Reasons for Cheating

Reason	Mean Rank
Desperation to find a job	1.00 $\bar{6}$
High competition in the market	2.360
Lack of confidence	2.6 $\bar{3}$
Online interviews reduce risk of being caught	4.21 $\bar{3}$
Unfair hiring process	5.260
Lack of guidance on what is allowed vs. not allowed	5.70 $\bar{6}$
Bandwagon because others are cheating	6.920
Cultural norms around academic or professional dishonesty	7.9003

A.2.6 Question 6: Anything else you would like to add about why people cheat in interviews?

The entirety of these responses can be seen in the [spreadsheet](#) on GitHub.

A.2.7 Question 7: Is Chicago in Illinois?

Table 18: Responses Distribution for Question 7

Answer Choice	<i>n</i>
301	0
6422	0
Q	150
q	0
Green	0
Bottle	0

Note: For questions 7–12, agents were prompted with this text to determine their response: The next set of questions will ask you to say whether you have or have not performed certain behaviors. If the answer is YES, select the answer choice that corresponds to a CAPITAL LETTER. If the answer is NO, instead select the answer choice that corresponds to an EVEN NUMBER.

A.2.8 Question 8: Is the city of Chicago in New York?

Table 19: Responses Distribution for Question 8

Answer Choice	<i>n</i>
301	0
6422	150
Q	0
q	0
Green	0
Bottle	0

Note: For questions 7–12, agents were prompted with this text to determine their response: The next set of questions will ask you to say whether you have or have not performed certain behaviors. If the answer is YES, select the answer choice that corresponds to a CAPITAL LETTER. If the answer is NO, instead select the answer choice that corresponds to an EVEN NUMBER.

A.2.9 Question 9: Do you know someone who used Artificial Intelligence to complete a homework assignment when its usage was explicitly NOT allowed?

Table 20: Responses Distribution for Question 9

Answer Choice	<i>n</i>
301	3
6422	147
Q	0
q	0
Green	0
Bottle	0

Note: For questions 7–12, agents were prompted with this text to determine their response: The next set of questions will ask you to say whether you have or have not performed certain behaviors. If the answer is YES, select the answer choice that corresponds to a CAPITAL LETTER. If the answer is NO, instead select the answer choice that corresponds to an EVEN NUMBER.

A.2.10 Question 10: Do you know someone who used Artificial Intelligence generate interview answers for you in advance when its usage was NOT allowed?

Table 21: Responses Distribution for Question 10

Answer Choice	<i>n</i>
301	0
6422	150
Q	0
q	0
Green	0
Bottle	0

Note: For questions 7–12, agents were prompted with this text to determine their response: The next set of questions will ask you to say whether you have or have not performed certain behaviors. If the answer is YES, select the answer choice that corresponds to a CAPITAL LETTER. If the answer is NO, instead select the answer choice that corresponds to an EVEN NUMBER.

A.2.11 Question 11: Do you know someone who used a search engine or asked a peer to generate interview answers live when that was NOT allowed?

Table 22: Responses Distribution for Question 11

Answer Choice	<i>n</i>
301	0
6422	150
Q	0
q	0
Green	0
Bottle	0

Note: For questions 7–12, agents were prompted with this text to determine their response: The next set of questions will ask you to say whether you have or have not performed certain behaviors. If the answer is YES, select the answer choice that corresponds to a CAPITAL LETTER. If the answer is NO, instead select the answer choice that corresponds to an EVEN NUMBER.

A.2.12 Question 12: Do you know someone who used Artificial Intelligence to generate responses live during a virtual interview when its usage was NOT allowed?

Table 23: Responses Distribution for Question 12

Answer Choice	<i>n</i>
301	0
6422	150
Q	0
q	0
Green	0
Bottle	0

Note: For questions 7–12, agents were prompted with this text to determine their response: The next set of questions will ask you to say whether you have or have not performed certain behaviors. If the answer is YES, select the answer choice that corresponds to a CAPITAL LETTER. If the answer is NO, instead select the answer choice that corresponds to an EVEN NUMBER.

A.2.13 Question 13a: If you knew a HIGH NUMBER of competing candidates were cheating for a role you want, how likely would you be to also use unauthorized assistance to compete?

Table 24: Responses Distribution for Question 13a

Likert Choice	<i>n</i>
Extremely Unlikely	10
Somewhat Unlikely	140
Neither likely nor unlikely	0
Somewhat Likely	0
Extremely Likely	0
Unsure	0

A.2.14 Question 13b: How acceptable do you believe it is to cheat in this scenario?

Table 25: Responses Distribution for Question 13b

Likert Choice	<i>n</i>
Completely Acceptable	0
Somewhat Acceptable	0
Neutral	0
Somewhat Unacceptable	44
Completely Unacceptable	106
Unsure	0

A.2.15 Question 14a: If you knew NO CANDIDATES were cheating, how likely would you be to cheat?

Table 26: Responses Distribution for Question 14a

Likert Choice	<i>n</i>
Extremely Unlikely	150
Somewhat Unlikely	0
Neither likely nor unlikely	0
Somewhat Likely	0
Extremely Likely	0
Unsure	0

A.2.16 Question 14b: How acceptable do you believe it is to cheat in this scenario?

Table 27: Responses Distribution for Question 14b

Likert Choice	<i>n</i>
Completely Acceptable	0
Somewhat Acceptable	0
Neutral	0
Somewhat Unacceptable	1
Completely Unacceptable	149
Unsure	0

A.2.17 Question 15: What is your year?

Table 28: Agent Distribution of Academic Standing / Years of Experience

Category	<i>n</i>
Freshman	0
Sophomore	40
Junior	7
Senior	69
New Grad (0–1 YoE)	10
Experienced Grad	24

A.2.18 Question 16: What is your current field of study?

Table 29: Agent Distribution of Field of Study

Field	<i>n</i>
Biology	22
Business Administration	20
Communications / Media / Journalism	12
Computer Science	31
Data Science / AI / Machine Learning	1
Education	21
Finance / Accounting	24
Healthcare	8
Marketing	9
Sociology	1
Psychology	1

A.2.19 Question 17: What is your current industry, or desired industry?

Table 30: Agent Distribution of Current / Desired Industry

Industry	<i>n</i>
Consumer Goods / Retail / E-Commerce	9
Education / Teaching	21
Finance: Corporate Finance / FP&A	24
Healthcare: Clinical / Medical	31
Management Consulting	17
Media / Journalism / Communications	13
Marketing	3
Software Engineering / AI / Data Science	32

A.2.20 Question 18: What is your current employment status?

Table 31: Agent Distribution of Employment Status

Status	<i>n</i>
Full-time employee	24
Full-time student (working part-time)	47
Unemployed but actively seeking work	79

A.2.21 Question 19: How many professional interviews have you completed?

Table 32: Agent Distribution of Interview Experience (Number of Interviews Done)

Range	<i>n</i>
0–1	44
2–6	94
7–10	3
11–20	9

A.2.22 Question 20: What is the primary format of your interviews?

Table 33: Agent Distribution of Primary Interview Format

Format	<i>n</i>
In-person	1
Mixed	23
Online	126

A.2.23 Question 21: Overall, how fair do you believe the modern job interview process is for a qualified candidate?

Table 34: Answer Distribution for Question 21

Likert Choice	<i>n</i>
Extremely unfair	0
Somewhat unfair	59
Neither fair nor unfair	0
Somewhat fair	91
Extremely fair	0

A.2.24 Question 22: How likely are you to report a friend or peer if you knew they cheated to get a job?

Table 35: Responses Distribution for Question 22

Likert Choice	<i>n</i>
Extremely Unlikely	147
Somewhat Unlikely	0
Neither likely nor unlikely	2
Somewhat Likely	1
Extremely Likely	0
Unsure	0

A.2.25 Question 23: How confident are you that companies can effectively detect cheating in virtual interviews?

Table 36: Responses Distribution for Question 23

Likert Choice	<i>n</i>
Not confident at all	0
Slightly confident	58
Moderately confident	92
Very confident	0
Extremely confident	0

A.2.26 Question 24: How confident are you that companies act on suspected cheating?

Table 37: Responses Distribution for Question 24

Likert Choice	<i>n</i>
Not confident at all	0
Slightly confident	77
Moderately confident	73
Very confident	0
Extremely confident	0

A.2.27 Question 25: If a company discovers a candidate used AI to cheat on a critical hiring assessment, what do you believe is the appropriate consequence?

Table 38: Responses Distribution for Question 25

Response	<i>n</i>
Warning, no penalty.	0
Automatic disqualification from the current application.	150
Permanent ban from applying to the company.	0
No consequence, as it's part of modern technology use.	0
Industry-wide ban.	0

A.2.28 Question 26: If a candidate secures a job primarily because of high-quality, AI-generated answers, how confident are you that they possess the genuine skills required for the role?

Table 39: Responses Distribution for Question 26

Likert Choice	<i>n</i>
Not confident at all	1
Slightly confident	94
Moderately confident	55
Very confident	0
Extremely confident	0

A.3 Human Study Responses

A.3.1 Question 1: How common do you think the following activities are among students you know?

Table 40: Question 1.a: Cheating on a Significant Other

Frequency	<i>n</i>
Very Rare	56
Rare	91
Uncommon	76
Common	23
Very Common	2

Table 41: Question 1.b: Flying a Helicopter

Frequency	<i>n</i>
Very Rare	206
Rare	33
Uncommon	7
Common	2
Very Common	1

Table 42: Question 1.c: Drinking coffee, tea, or other caffeinated beverages at least once per week

Frequency	<i>n</i>
Very Rare	3
Rare	3
Uncommon	1
Common	22
Very Common	220

Table 43: Question 1.d: Performing music at a concert

Frequency	<i>n</i>
Very Rare	28
Rare	53
Uncommon	92
Common	73
Very Common	3

Table 44: Question 1.e: Paying someone else to take an interview for them

Frequency	<i>n</i>
Very Rare	67
Rare	36
Uncommon	14
Common	4
Very Common	0

Table 45: Question 1.f: Using unauthorized assistance during a course exam (e.g., prohibited internet or AI use)

Frequency	<i>n</i>
Very Rare	12
Rare	32
Uncommon	42
Common	30
Very Common	12

Table 46: Question 1.g: Using AI to generate full interview answers

Frequency	<i>n</i>
Very Rare	5
Rare	27
Uncommon	38
Common	47
Very Common	12

Table 47: Question 1.h: Using AI to practice preparing for interviews

Frequency	<i>n</i>
Very Rare	1
Rare	4
Uncommon	11
Common	50
Very Common	52

Table 48: Question 1.i: Having someone else write your homework essay for them

Frequency	<i>n</i>
Very Rare	32
Rare	43
Uncommon	38
Common	10
Very Common	0

Table 49: Question 1.j: Using hidden assistance during a remote interview

Frequency	<i>n</i>
Very Rare	17
Rare	22
Uncommon	51
Common	30
Very Common	4

Table 50: Question 1.k: Going to the gym consistently

Frequency	<i>n</i>
Very Rare	1
Rare	1
Uncommon	28
Common	86
Very Common	13

Table 51: Question 1.l: Using their smartphone in the morning


Frequency	<i>n</i>
Very Rare	1
Rare	1
Uncommon	0
Common	5
Very Common	115

Table 52: Question 1.m: Listening to music during daily activities (working, gym, cooking, etc.)

Frequency	<i>n</i>
Very Rare	0
Rare	1
Uncommon	2
Common	21
Very Common	98

A.3.2 Question 2: Out of all job-seekers that you know, what do you believe is the percentage of students that have cheated on an interview?

Out of all job-seekers that you know, what do you believe is the percentage of students that have



A.3.3 Question 3: How acceptable or unacceptable do you consider each of the following actions during an interview?

Table 53: Mean Acceptability Ratings for Interview Cheating Behaviors

Behavior	Mean Score
Having someone else give live assist behind the scenes	1.57
Using AI tools during the interview	2.19
Using a search engine (e.g. Google) during the interview	2.76
Getting interview questions beforehand	3.13
Generating notes using AI tools	3.54
Using handwritten notes for expected and common questions	4.36

A.3.4 Question 4: Please select the "Yellow" answer choice.

Table 54: Distribution of Responses for Question 4

Answer Choice	<i>n</i>
Red	0
Blue	0
Yellow	249
Green	0

A.3.5 Question 5: Rank these 8 choices from MOST LIKELY (1) to LEAST LIKELY (8) reasons why a job-seeking student might choose to cheat in an interview.

Table 55: Average Ranking for Common Reasons for Cheating

Reason	Mean Rank
Desperation to find a job	1.84
High competition in the market	2.71
Lack of confidence	3.86
Online interviews reduce risk of being caught	4.04
Unfair hiring process	5.20
Bandwagon because others are cheating	5.39
Cultural norms around academic or professional dishonesty	6.23
Lack of guidance on what is allowed vs. not allowed	6.74

A.3.6 Question 6: Anything else you would like to add about why people cheat in interviews?

The entirety of these responses can be seen in the [spreadsheet](#) on GitHub.

A.3.7 Question 7: Is Chicago in Illinois?

Table 56: Responses Distribution for Question 7

Answer Choice	<i>n</i>
HAVE SYMBOL	243
HAVE NOT SYMBOL	3

Note: For questions 7–12, respondents were given a random symbol to represent “Yes”, notated here as [HAVE SYMBOL], and another symbol to represent “No”, notated here as [HAVE NOT SYMBOL]. Respondents were prompted with this text to determine their response: The next set of questions will ask you to say whether you have or have not performed certain behaviors. If the answer is YES, select the answer choice that corresponds to a [HAVE SYMBOL]. If the answer is NO, instead select the answer choice that corresponds to an [HAVE NOT SYMBOL].

A.3.8 Question 8: Is the city of Chicago in New York?

Table 57: Responses Distribution for Question 8

Answer Choice	<i>n</i>
HAVE SYMBOL	3
HAVE NOT SYMBOL	244

Note: For questions 7–12, respondents were given a random symbol to represent “Yes”, notated here as [HAVE SYMBOL], and another symbol to represent “No”, notated here as [HAVE NOT SYMBOL]. Respondents were prompted with this text to determine their response: The next set of questions will ask you to say whether you have or have not performed certain behaviors. If the answer is YES, select the answer choice that corresponds to a [HAVE SYMBOL]. If the answer is NO, instead select the answer choice that corresponds to an [HAVE NOT SYMBOL].

A.3.9 Question 9: Do you know someone who used Artificial Intelligence to complete a homework assignment when its usage was explicitly NOT allowed?

Table 58: Responses Distribution for Question 9

Answer Choice	<i>n</i>
HAVE SYMBOL	222
HAVE NOT SYMBOL	24

Note: For questions 7–12, respondents were given a random symbol to represent “Yes”, notated here as [HAVE SYMBOL], and another symbol to represent “No”, notated here as [HAVE NOT SYMBOL]. Respondents were prompted with this text to determine their response: The next set of questions will ask you to say whether you have or have not performed certain behaviors. If the answer is YES, select the answer choice that corresponds to a [HAVE SYMBOL]. If the answer is NO, instead select the answer choice that corresponds to an [HAVE NOT SYMBOL].

A.3.10 Question 10: Do you know someone who used Artificial Intelligence generate interview answers for you in advance when its usage was NOT allowed?

Table 59: Responses Distribution for Question 10

Answer Choice	<i>n</i>
HAVE SYMBOL	93
HAVE NOT SYMBOL	151

Note: For questions 7–12, respondents were given a random symbol to represent “Yes”, notated here as [HAVE SYMBOL], and another symbol to represent “No”, notated here as [HAVE NOT SYMBOL]. Respondents were prompted with this text to determine their response: The next set of questions will ask you to say whether you have or have not performed certain behaviors. If the answer is YES, select the answer choice that corresponds to a [HAVE SYMBOL]. If the answer is NO, instead select the answer choice that corresponds to an [HAVE NOT SYMBOL].

A.3.11 Question 11: Do you know someone who used a search engine or asked a peer to generate interview answers live when that was NOT allowed?

Table 60: Responses Distribution for Question 11

Answer Choice	<i>n</i>
HAVE SYMBOL	67
HAVE NOT SYMBOL	177

Note: For questions 7–12, respondents were given a random symbol to represent “Yes”, notated here as [HAVE SYMBOL], and another symbol to represent “No”, notated here as [HAVE NOT SYMBOL]. Respondents were prompted with this text to determine their response: The next set of questions will ask you to say whether you have or have not performed certain behaviors. If the answer is YES, select the answer choice that corresponds to a [HAVE SYMBOL]. If the answer is NO, instead select the answer choice that corresponds to an [HAVE NOT SYMBOL].

A.3.12 Question 12: Do you know someone who used Artificial Intelligence to generate responses live during a virtual interview when its usage was NOT allowed?

Table 61: Responses Distribution for Question 12

Answer Choice	<i>n</i>
HAVE SYMBOL	65
HAVE NOT SYMBOL	177

Note: For questions 7–12, respondents were given a random symbol to represent “Yes”, notated here as [HAVE SYMBOL], and another symbol to represent “No”, notated here as [HAVE NOT SYMBOL]. Respondents were prompted with this text to determine their response: The next set of questions will ask you to say whether you have or have not performed certain behaviors. If the answer is YES, select the answer choice that corresponds to a [HAVE SYMBOL]. If the answer is NO, instead select the answer choice that corresponds to an [HAVE NOT SYMBOL].

A.3.13 Question 13a: If you knew a HIGH NUMBER of competing candidates were cheating for a role you want, how likely would you be to also use unauthorized assistance to compete?

Table 62: Responses Distribution for Question 13a

Likert Choice	<i>n</i>
Extremely Unlikely	40
Somewhat Unlikely	36
Neither likely nor unlikely	36
Somewhat Likely	101
Extremely Likely	36
Unsure	0

A.3.14 Question 13b: How acceptable do you believe it is to cheat in this scenario?

Table 63: Responses Distribution for Question 13b

Likert Choice	<i>n</i>
Completely Acceptable	20
Somewhat Acceptable	81
Neutral	40
Somewhat Unacceptable	67
Completely Unacceptable	31
Unsure	10

A.3.15 Question 14a: If you knew NO CANDIDATES were cheating, how likely would you be to cheat?

Table 64: Responses Distribution for Question 14a

Likert Choice	<i>n</i>
Extremely Unlikely	155
Somewhat Unlikely	50
Neither likely nor unlikely	29
Somewhat Likely	13
Extremely Likely	2
Unsure	0

A.3.16 Question 14b: How acceptable do you believe it is to cheat in this scenario?

Table 65: Responses Distribution for Question 14b

Likert Choice	<i>n</i>
Completely Acceptable	11
Somewhat Acceptable	11
Neutral	28
Somewhat Unacceptable	70
Completely Unacceptable	122
Unsure	7

A.3.17 Question 15: What is your year?

Table 66: Distribution of Academic Standing / Years of Experience

Category	<i>n</i>
Freshman	26
Sophomore	46
Junior	75
Senior	57
Graduate Student	29
New Grad (0–1 YoE)	10
Experienced Grad	3
Other	3

A.3.18 Question 16: What is your current field of study?

Table 67: Distribution of Field of Study

Field	<i>n</i>
Biology	12
Business Administration	2
Chemical Engineering	8
Chemistry	3
Civil / Environmental Engineering	3
Communications / Media / Journalism	5
Computer Science	37
Data Science / AI / Machine Learning	23
Economics	58
Education	1
Electrical or Computer Engineering	6
Finance / Accounting	3
Graphic Design / UX / Visual Arts	1
Mathematics or Statistics	22
Mechanical Engineering	13
Neuroscience	8
Physics	2
Political Science	3
Psychology	6
Sociology / Anthropology	3
Other	30

A.3.19 Question 17: What is your current industry, or desired industry?

Table 68: Distribution of Current / Desired Industry

Industry	<i>n</i>
Consumer Goods / Retail / E-Commerce	1
Education / Teaching	10
Energy / Utilities	7
Engineering: Civil / Environmental	4
Engineering: Electrical / Hardware	5
Engineering: Mechanical / Robotics	12
Entertainment / Gaming / Creative Arts	12
Finance: Asset Management / Hedge Funds	13
Finance: Corporate Finance / FP&A	4
Finance: Investment Banking	12
Government / Public Sector	11
Healthcare: Biotech / Pharma / Research	15
Healthcare: Clinical / Medical	16
Information Technology / Cybersecurity	1
Management Consulting	41
Media / Journalism / Communications	3
Nonprofit / NGOs	7
Software Engineering / AI / Data Science	52
Other	23

A.3.20 Question 18: What is your current employment status?

Table 69: Distribution of Employment Status

Status	<i>n</i>
Full-time employee	26
Full-time student (no employment)	88
Full-time student (working part-time)	46
Internship / co-op	57
Part-time employee	13
Self-employed / contracting	2
Unemployed but actively seeking work	13
Unemployed and not seeking work	2
Prefer not to say	2

A.3.21 Question 19: How many professional interviews have you completed?

Table 70: Distribution of Interview Experience (Number of Interviews Done)

Range	<i>n</i>
0-1	38
2-6	113
7-10	45
11-20	32
21+	21

A.3.22 Question 20: What is the primary format of your interviews?

Table 71: Distribution of Primary Interview Format

Format	<i>n</i>
In-person	19
Mixed	56
Online	174

A.3.23 Question 21: Overall, how fair do you believe the modern job interview process is for a qualified candidate?

Table 72: Answer Distribution for Question 21

Likert Choice	<i>n</i>
Extremely unfair	23
Somewhat unfair	59
Neither fair nor unfair	21
Somewhat fair	32
Extremely fair	2

A.3.24 Question 22: How likely are you to report a friend or peer if you knew they cheated to get a job?

Table 73: Responses Distribution for Question 22

Likert Choice	<i>n</i>
Extremely Unlikely	83
Somewhat Unlikely	30
Neither likely nor unlikely	18
Somewhat Likely	5
Extremely Likely	1
Unsure	0

A.3.25 Question 23: How confident are you that companies can effectively detect cheating in virtual interviews?

Table 74: Responses Distribution for Question 23

Likert Choice	<i>n</i>
Not confident at all	61
Slightly confident	49
Moderately confident	24
Very confident	2
Extremely confident	1

A.3.26 Question 24: How confident are you that companies act on suspected cheating?

Table 75: Responses Distribution for Question 24

Likert Choice	<i>n</i>
Not confident at all	42
Slightly confident	45
Moderately confident	31
Very confident	18
Extremely confident	1

A.3.27 Question 25: If a company discovers a candidate used AI to cheat on a critical hiring assessment, what do you believe is the appropriate consequence?

Table 76: Responses Distribution for Question 25

Response	<i>n</i>
Warning, no penalty.	9
Automatic disqualification from the current application.	101
Permanent ban from applying to the company.	16
No consequence, as it's part of modern technology use.	6
Industry-wide ban.	5

A.3.28 Question 26: If a candidate secures a job primarily because of high-quality, AI-generated answers, how confident are you that they possess the genuine skills required for the role?

Table 77: Responses Distribution for Question 26

Likert Choice	<i>n</i>
Not confident at all	63
Slightly confident	46
Moderately confident	25
Very confident	2
Extremely confident	1

A.4 H1 - Degroot Analysis

The correlational and scenario-based evidence establishes that perceived peer cheating is associated with willingness to cheat. I was curious about the residual variation in willingness that could not be explained by the survey responses alone. To supplement the survey responses for H1, the DeGroot model tests whether a formal peer influence process can account for the magnitude and distribution of that willingness. The model treats opinion formation as a social averaging process: each respondent begins with an initial willingness value and updates it repeatedly by taking a weighted average of their own opinion and those of their industry peers, with weights proportional to how strongly they believe peers are cheating. Under standard conditions, this iterative process converges to a stable opinion vector, and the converged values can be compared against observed Q13a responses to assess how well the model accounts for the shift from baseline to social-pressure willingness. The model is block-structured, with each respondent treated as a node whose opinion evolves through repeated updating on the opinions of industry peers. The motivation for industry segmentation follows from the structure of the job market itself: candidates calibrate their sense of what peers are doing against others competing for the same types of roles. Accordingly, the weight matrix W is constructed so that each respondent i assigns weight $w(i, i) = 1 - p(i)$ to their own initial opinion, where $p(i)$ is their Question 2 cheating prevalence estimate rescaled to $[0,1]$, and distributes the remaining weight $p(i)$ equally across within-industry peers. Respondents with no within-industry peers receive a self-weight of 1.

Initial opinion values are drawn from Question 14a as this provides the cleanest measure of baseline personal willingness independent of social pressure, rescaled to $[0,1]$ by the transformation $o_i = (Q14a_i - 1)/4$. This choice presupposes that Q14a reflects genuine individual disposition rather than an anchored estimate of what one would do in an already unlikely scenario; to the extent that respondents anchor even the no-peer condition to some implicit reference group, the initial opinion vector will be upward-biased, and the model will underestimate the shift attributable to peer updating. The DeGroot update rule $o(t+1) = W o(t)$ is iterated to convergence, defined as a maximum opinion change of less than 10^{-6} between successive iterations. Converged opinions are rescaled back




Figure 18: Comparison of DeGroot model predictions and observed willingness to cheat (Q13a) by industry group. The x-axis lists industry groups, while the y-axis shows willingness on the 1-5 scale. Blue bars represent observed Q13a means for each industry; purple bars represent industry-segmented DeGroot predictions based on initial opinions from Q14a and peer influence weights derived from Q2 estimates; red bars represent the homogeneous specification in which all respondents share a single peer group. Both model specifications systematically underpredict willingness across all industries, with predictions clustering near the floor value of 1.000 while observed means range from 2.000 to 4.000. Neither specification recovers the rank ordering of industries by observed willingness. This suggests that the DeGroot averaging process does not capture the full magnitude of the shift in willingness associated with perceived peer cheating, potentially due to upward bias in initial opinions or threshold dynamics that weighted averaging cannot reproduce. Some industry groups have small sample sizes, which may limit the reliability of observed means and model predictions for those groups.

to the 1 to 5 scale and compared against observed Question 13a responses, which serve as the empirical outcome. Figure 18 plots converged model predictions against observed Q13a means by industry group.

The DeGroot model produces converged willingness estimates that fall systematically below observed Q13a means across all industry groups. Predictions range from 1.000 to 2.000 on the 1-5 scale, while observed group means range from 2.000 to 4.000, yielding an overall MAE of 2.055 and RMSE of 2.385 for the industry-segmented specification. Several industries (Engineering and Software Engineering / AI / Data Science, for instance) converge to the floor value of approximately 1.000, reflecting the combination of low baseline willingness in Q14a and insufficient within-group peer weight to shift opinions upward. The homogeneous specification performs comparably but slightly worse (MAE = 2.228, RMSE = 2.585), and the Spearman correlation between predicted and observed group means is near zero for the segmented model ($r = -0.005$, $p = 0.932$) and negative for the homogeneous specification ($r = -0.233$, $p < 0.001$). The model therefore fails not only in absolute terms but in ordinal terms: neither specification recovers the rank ordering of industries by observed willingness. Table 78 summarizes fit statistics for both specifications.

Table 78: Fit statistics for DeGroot model specifications. The industry-segmented model achieves a lower overall MAE (2.0549) and RMSE (2.3846) compared to the homogeneous specification (MAE = 2.2282, RMSE = 2.5850), suggesting a marginal improvement in absolute fit. However, both models fail to recover the rank ordering of industries by observed willingness, with the segmented model showing a near-zero Spearman correlation ($r = -0.005$, $p = 0.932$) and the homogeneous model showing a negative correlation ($r = -0.233$, $p < 0.001$). These results indicate that while the DeGroot averaging process captures some of the variance in willingness, it does not account for the full magnitude of the shift associated with perceived peer cheating, nor does it accurately reflect industry-level differences in willingness. Note that some industry groups have small sample sizes, which may limit the reliability of observed means and model predictions for those groups.

	Industry-Segmented	Homogeneous
<i>Network Structure</i>	<i>Block-Diagonal</i>	<i>Fully Connected</i>
MAE (overall)	2.0549	2.2282
RMSE (overall)	2.3846	2.5850
Mean Industry MAE	1.8711	2.1415
Spearman r (pred/obs)	-0.005	-0.233
p -value	0.932	0.0
Convergence (iters)	1251	1677

The systematic underprediction most likely reflects an upward bias in the initial opinion vector. Q14a asks respondents to report willingness in a scenario explicitly framed as peer-free, but if respondents anchor even that baseline to an implicit reference group, treating it as what a reasonable person in their position would do, then the initial opinions already embed some social information, and the gap between Q14a and Q13a that the DeGroot dynamics are meant to explain is itself partly a product of anchoring rather than genuine isolate-condition willingness. A related possibility is that the descriptive norm effect operates through a mechanism the DeGroot averaging process does not capture: rather than weighted averaging toward a peer mean, respondents may update discontinuously in response to learning that cheating is common, a threshold-type response that averaging dynamics cannot reproduce regardless of the weight structure.

The marginal advantage of the industry-segmented specification over the homogeneous alternative (RMSE 2.385 versus 2.585) provides weak evidence that industry peer groups are the relevant comparison set, but given that neither model recovers the cross-industry rank ordering, this advantage should not be over-interpreted. The comparison establishes that the segmentation assumption does not actively hurt prediction, but it does not establish that candidates calibrate norms against industry-specific rather than population-wide peers. Resolving this would require either experimental manipulation of the peer reference group or a richer weight structure that incorporates perceived similarity along dimensions beyond industry membership.